



Ciências  
ULisboa

# **Programação em Sistemas Distribuídos**

**MEI-MI-MSI**

**2018/19**

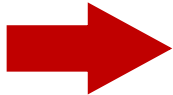
## **1. Review of Fundamental Distributed Systems Concepts**

**Prof. António Casimiro**

# What is a distributed system?

- A computer network *is not* a distributed system
- Computer network
  - infrastructure serving a set of computers interconnected through communication links of possibly diverse media and topology, and using a common set of communication protocols
- Distributed system
  - system composed of several computers which communicate through a **computer network**, hosting processes that use a common set of distributed protocols to assist the coherent execution of distributed activities

# Characteristics of Distributed Systems

- multiple computers
  - interconnected by a network
  - sharing state
- 
- independent failures
  - communication is unreliable
  - has variable delays
  - speed and bandwidth are moderate
  - investment costs often lower than mainframes
  - management costs are higher
  - partial ordering of events only
  - difficult to assess global state

# Centralized versus Distributed Systems



Ciências  
ULisboa

- Centralized :
  - Accessibility
    - to resources and info
  - Homogeneity
    - of procedures and technologies
  - Manageability
    - single, domain
  - Consistency
    - global state
  - Security, due to:
    - threat reduction
- Distributed :
  - Scalability, helped by:
    - geographical scope
    - heterogeneity
    - modularity
  - Sharing:
    - of common resources and info
  - Reliability and availability, due to:
    - redundancy and replication
    - graceful degradation
    - failure independence
  - Security, due to:
    - vulnerability reduction
  - Low cost factor



**Ciências**  
**ULisboa**

# **Distributed Systems Evolution**

## **TECHNIQUES AND ARCHITECTURES**

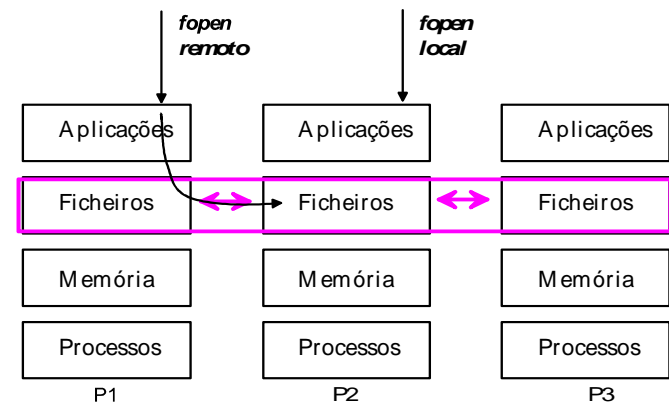
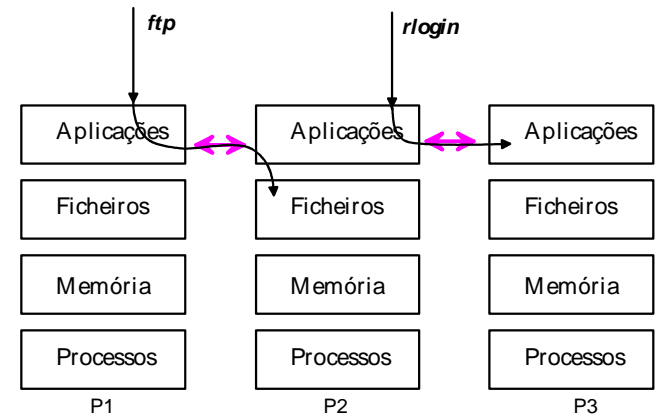
# Distributed Systems Evolution

## Techniques



Ciências  
ULisboa

- Resource sharing:
  - Remote login (rlogin)
  - File transfer
- Distributed file system:
  - Avoiding explicit file transfer
  - Efficient information sharing

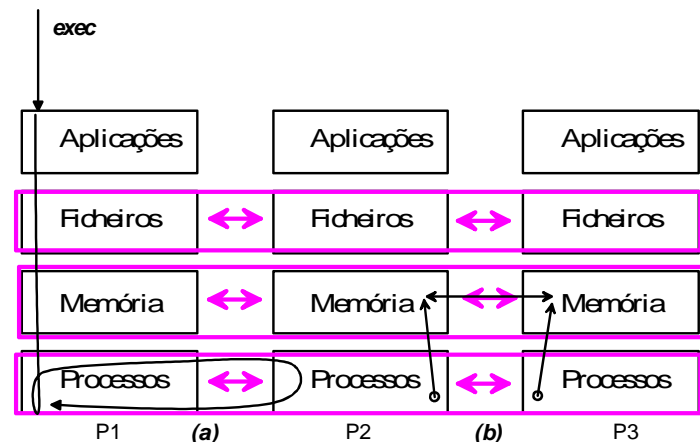
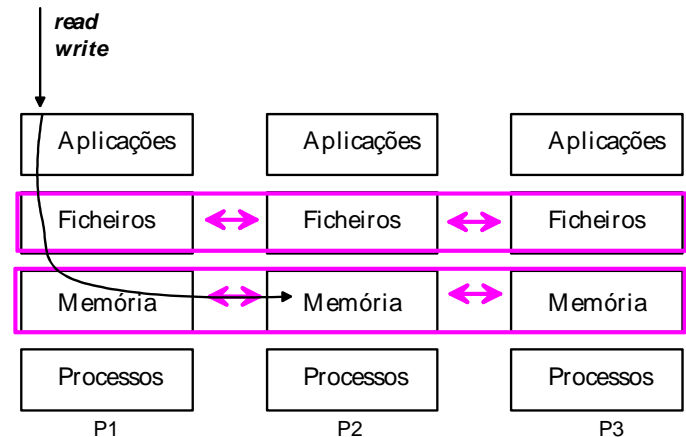


# Distributed Systems Evolution Techniques



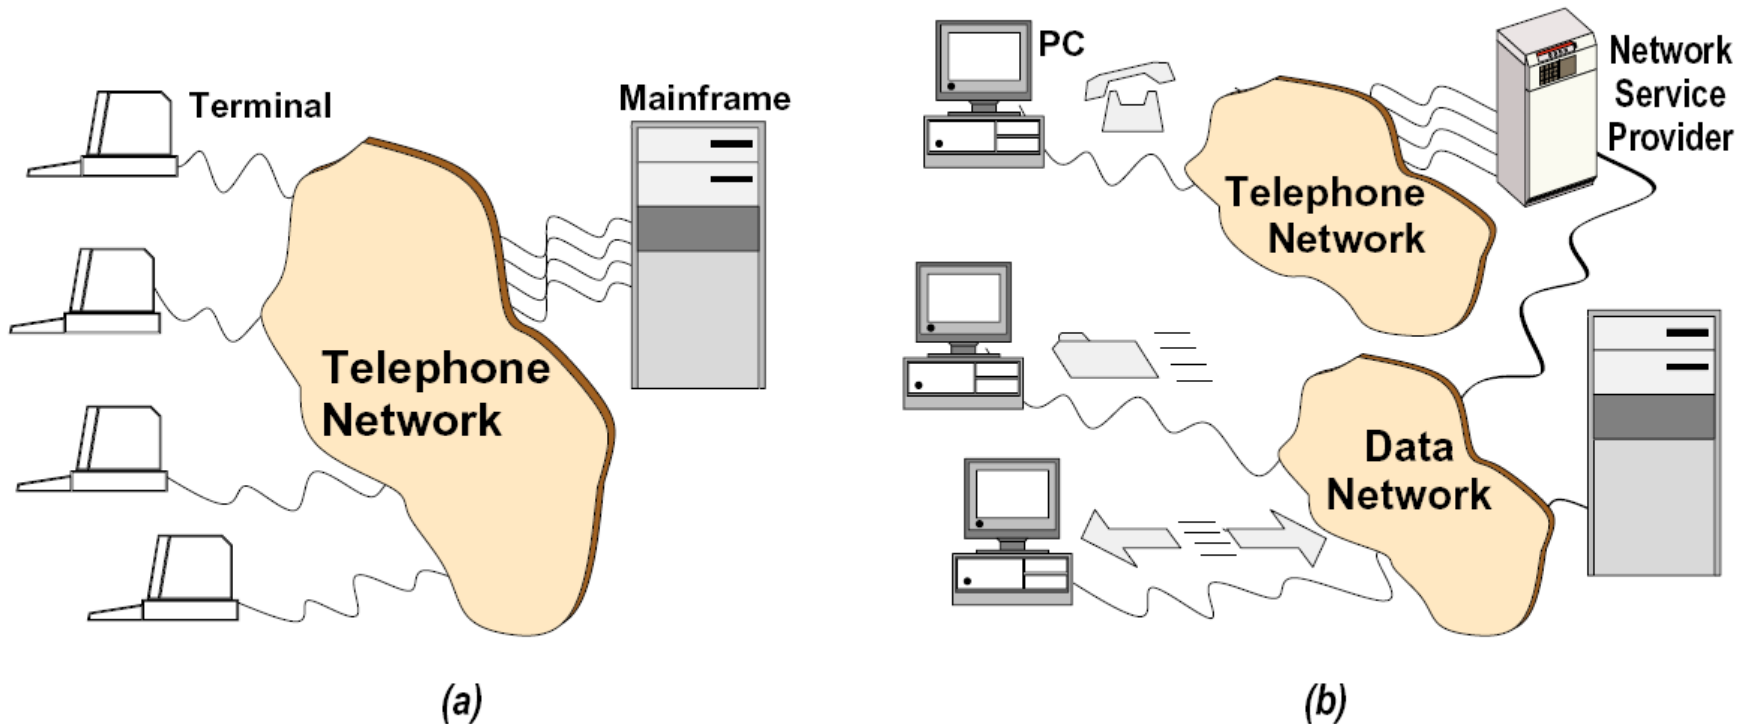
Ciências  
ULisboa

- Distributed memory:
  - Exploit network resources
  - Support (distributed) shared memory
  - Step towards distributed processing
- Distributed processing:
  - (a) Activities executed in several machines
  - (b) Concurrent activities synchronized over shared memory



# Remote Access Architectures:

## (a) Plain Telephone Line; (b) Data Network

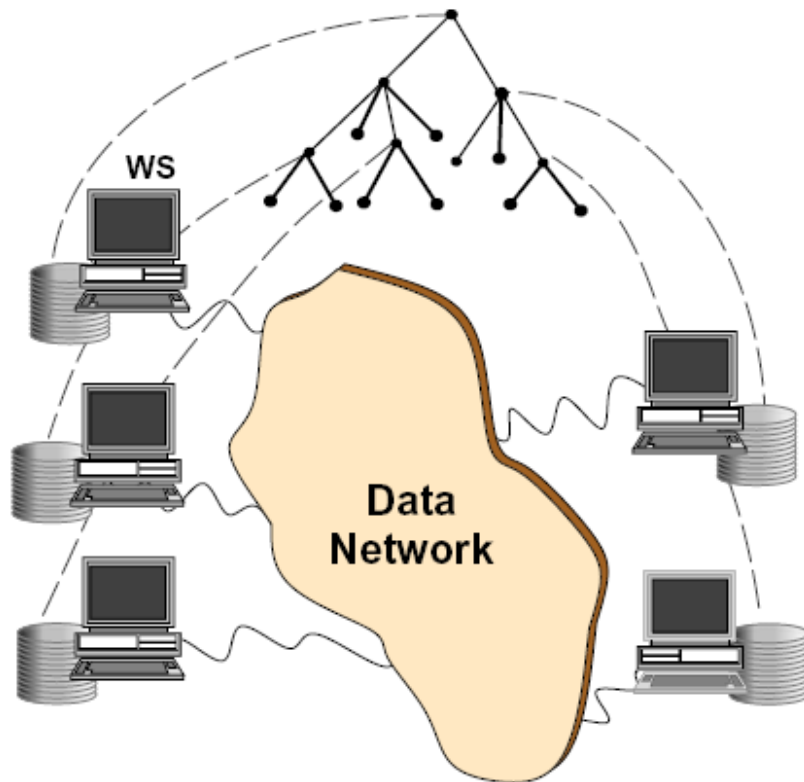




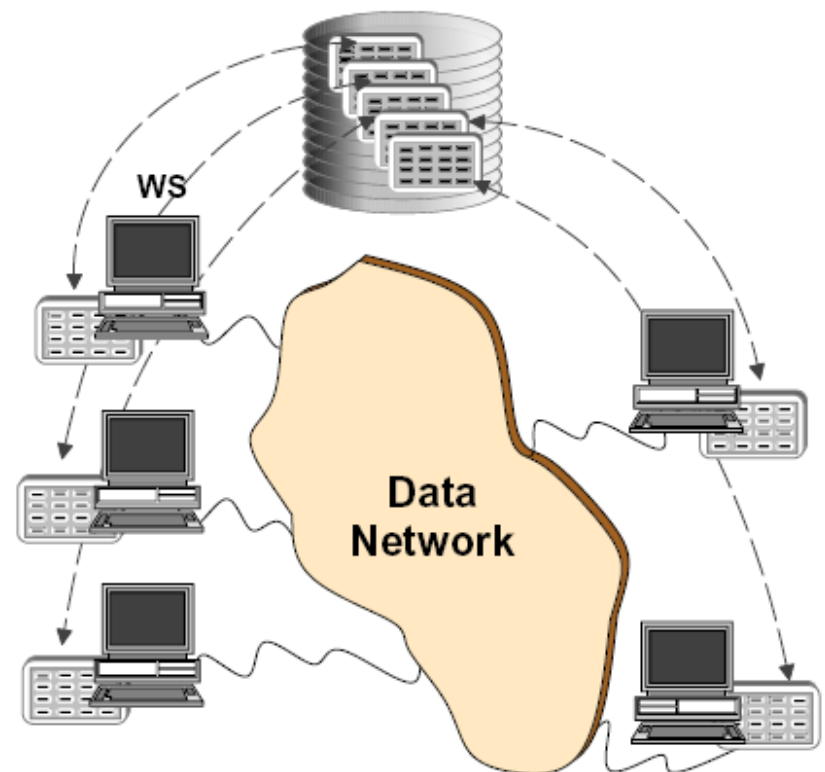
# Distributed File and Memory Architect.



Ciências  
ULisboa



(a)



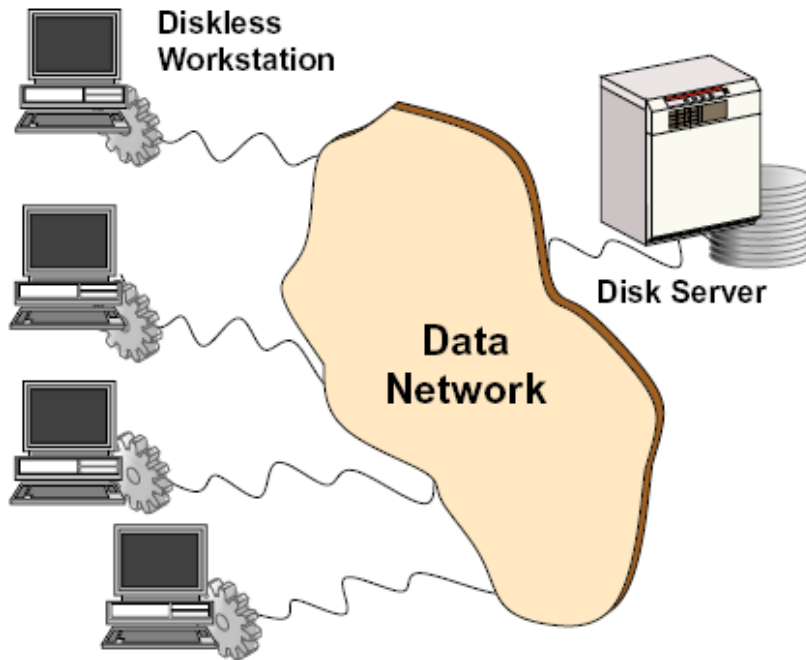
(b)

# Diskless and X-Terminal Architect.

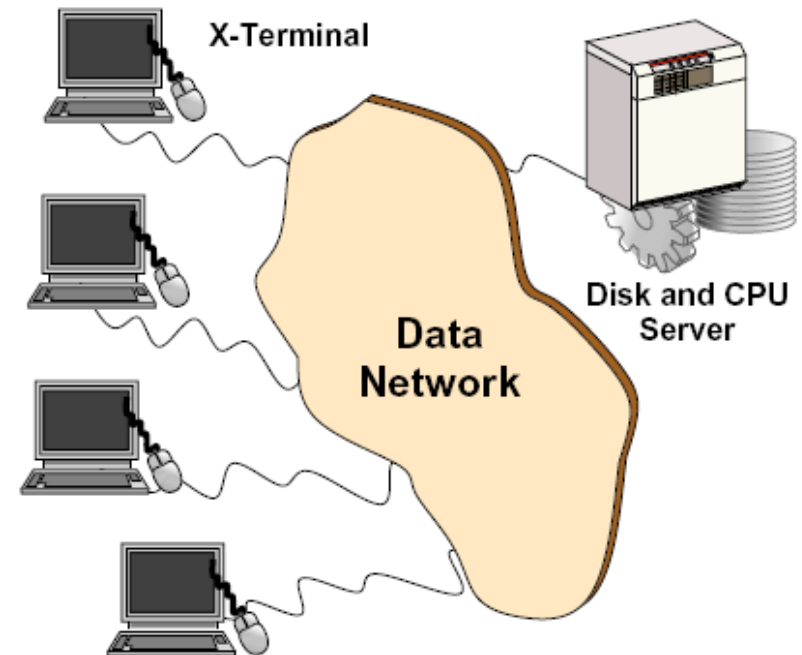
(on disk and/or CPU servers)



Ciências  
ULisboa



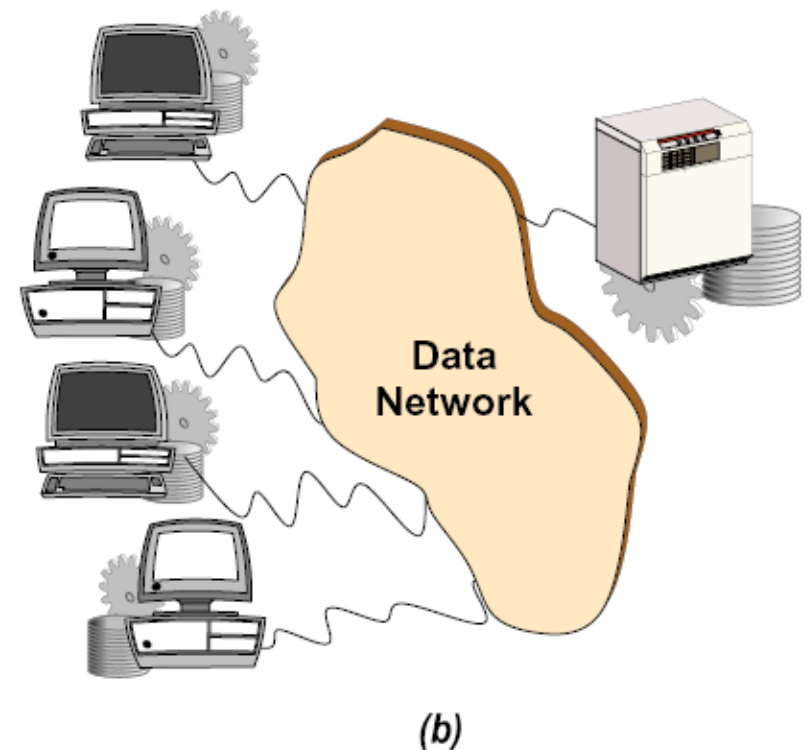
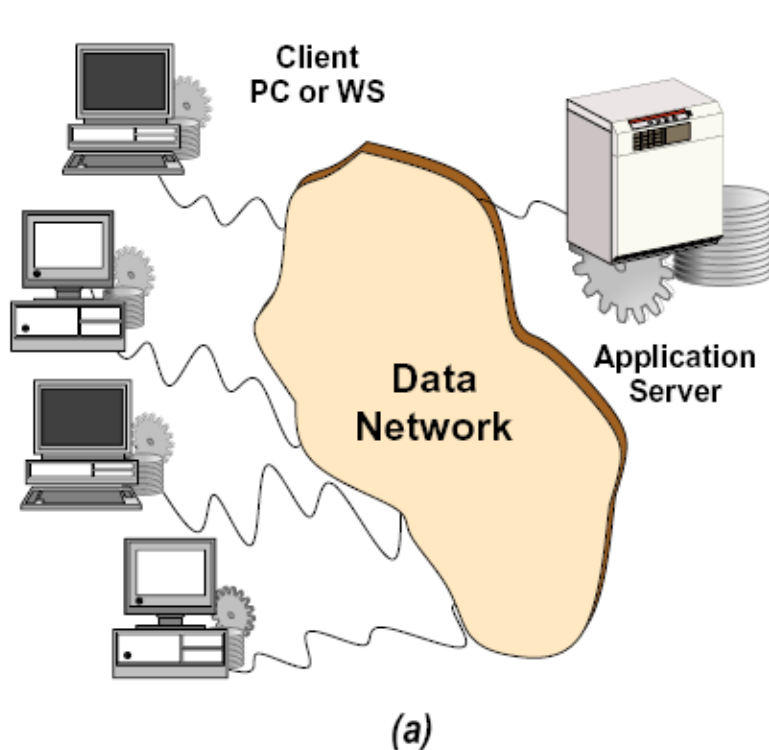
(a)



(b)

# Client-Server Architectures

(Clients became “fat”, perhaps too fat...)

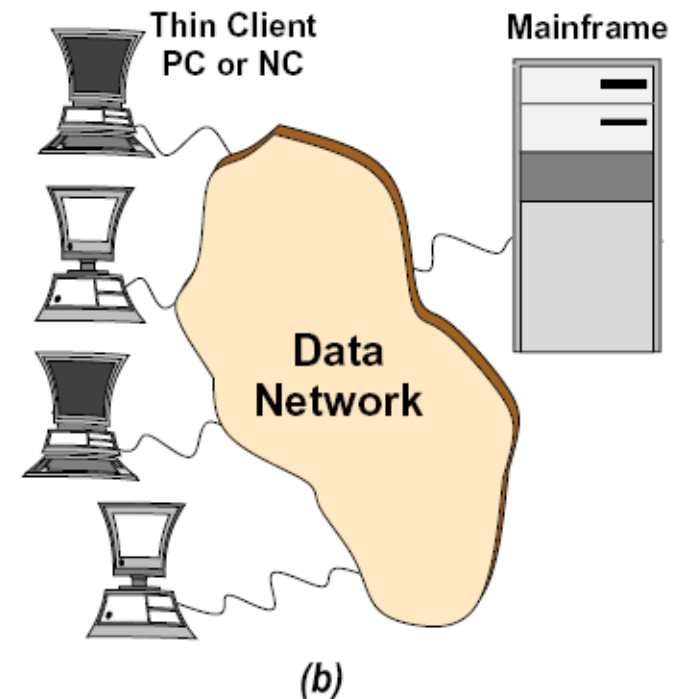
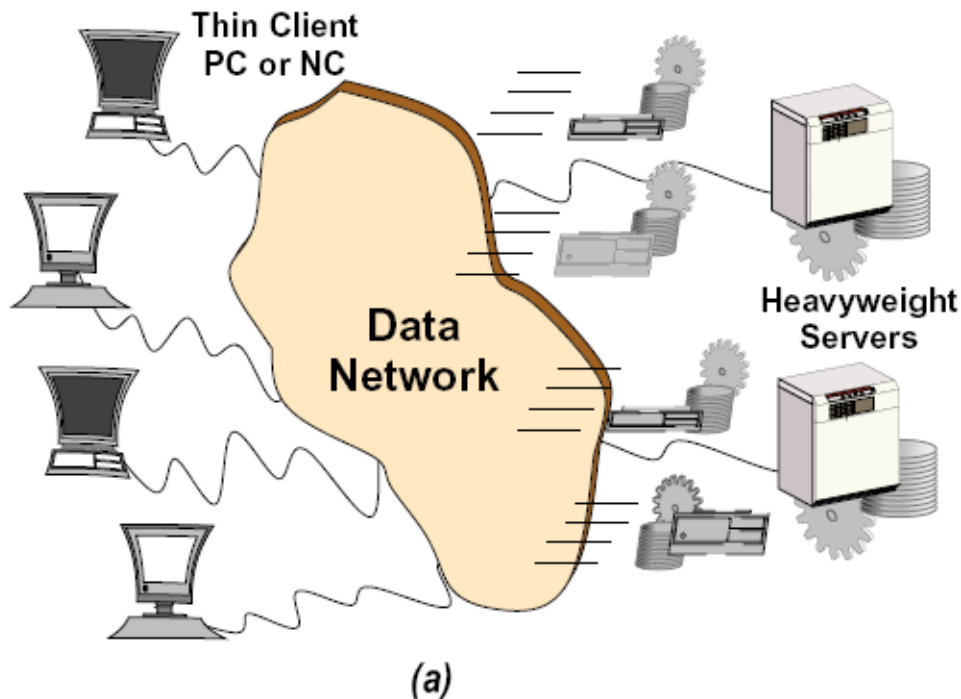


# 3-tier Client-Server or Thin-Client Arch.

(brought mainframes back, perhaps unnecessarily)



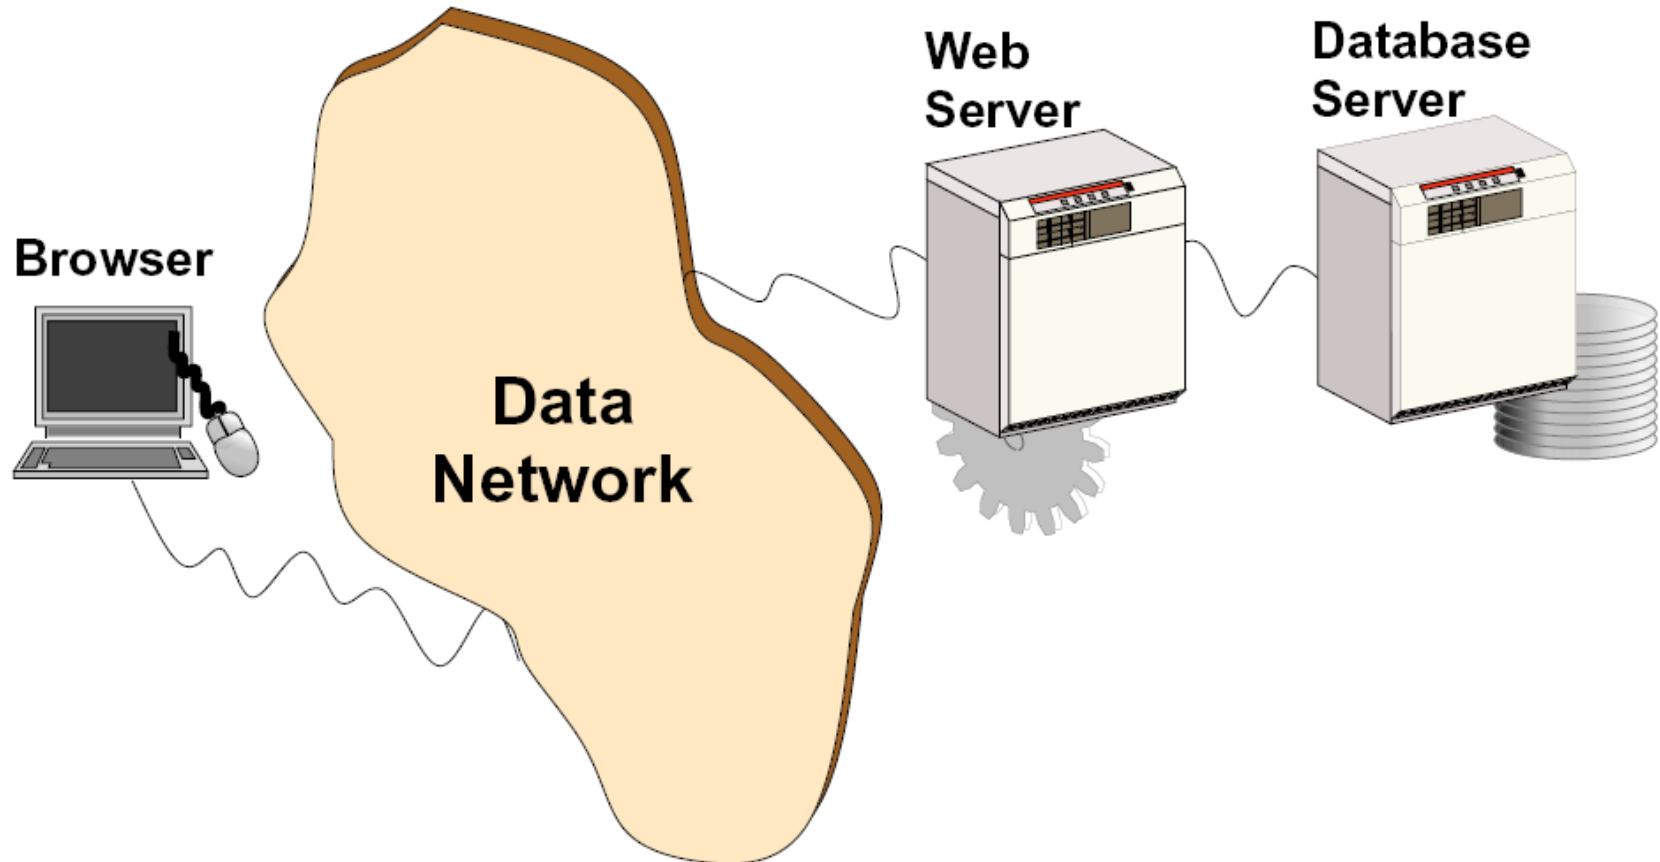
Ciências  
ULisboa



# 3-tier Client-Server Web-based Archtit.



Ciências  
ULisboa

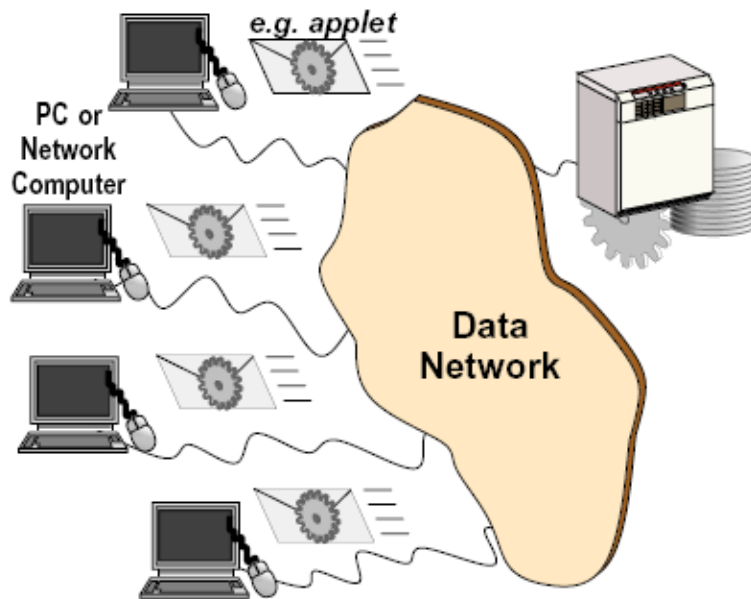


# Mobile Code Architectures:

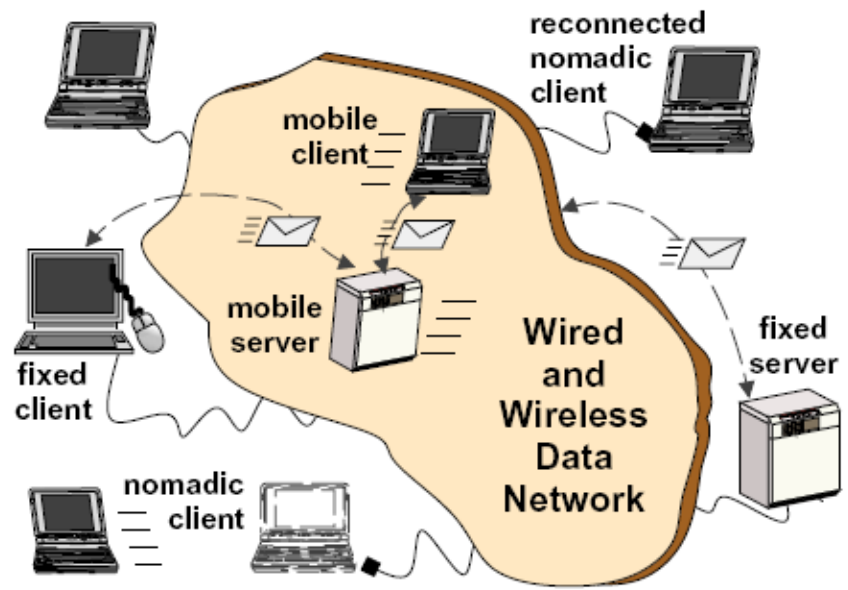
## (a) Portable and Mobile Code; (b) Mobile Nodes



Ciências  
ULisboa



(a)



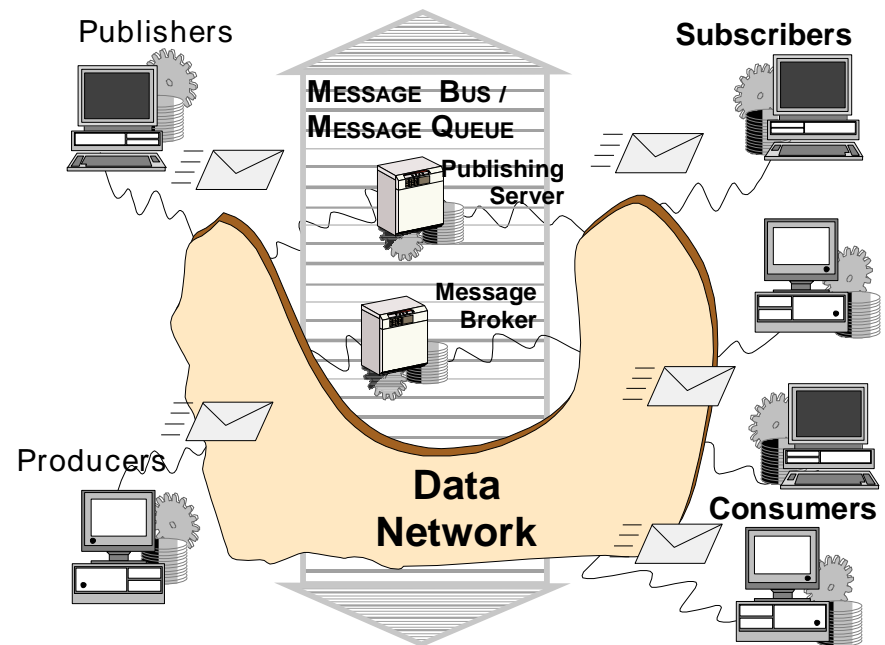
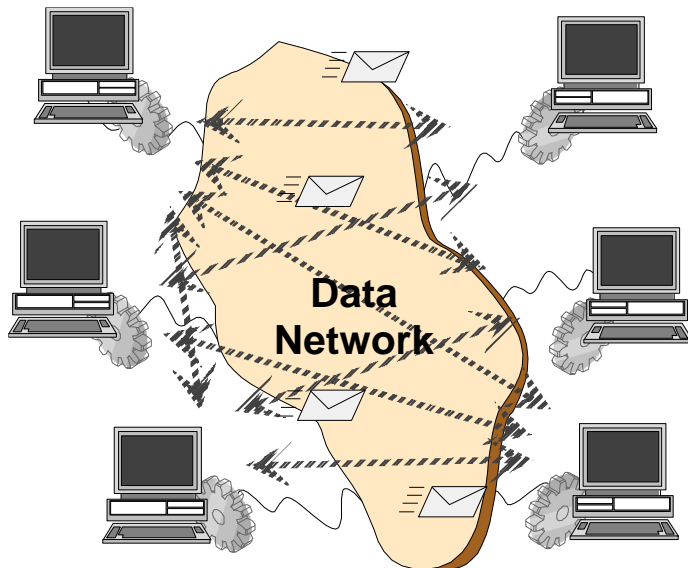
(b)

# Message and Event-based Architectures:

(a) Multipeer; (b) Publisher-subscriber, Message Queue



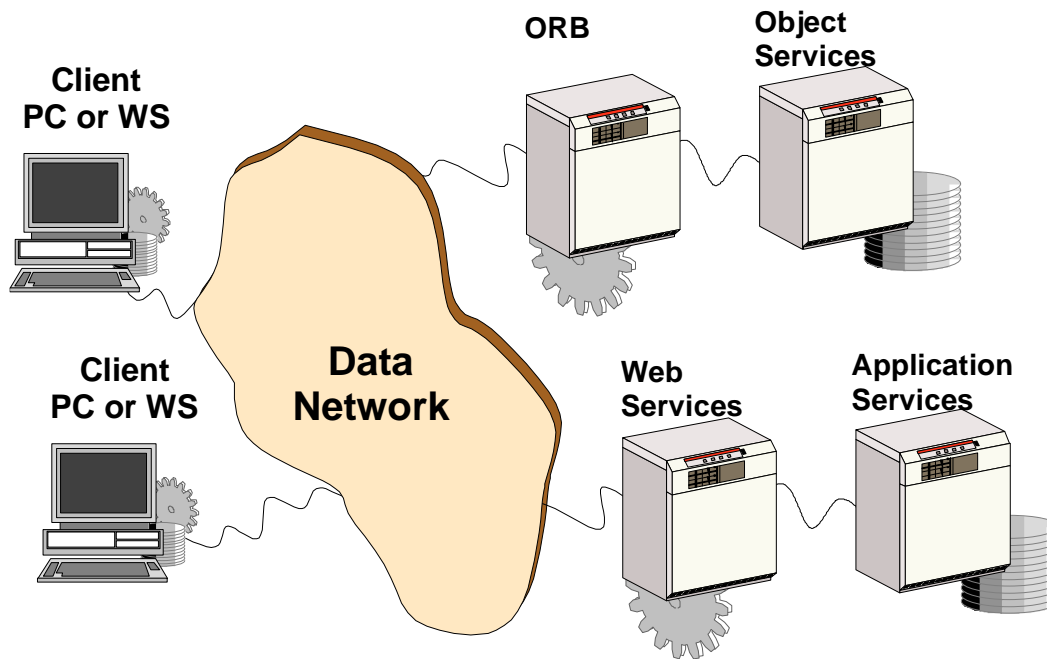
Ciências  
ULisboa



# Distributed Objects, Web Services (Web-based Architectures episode II)

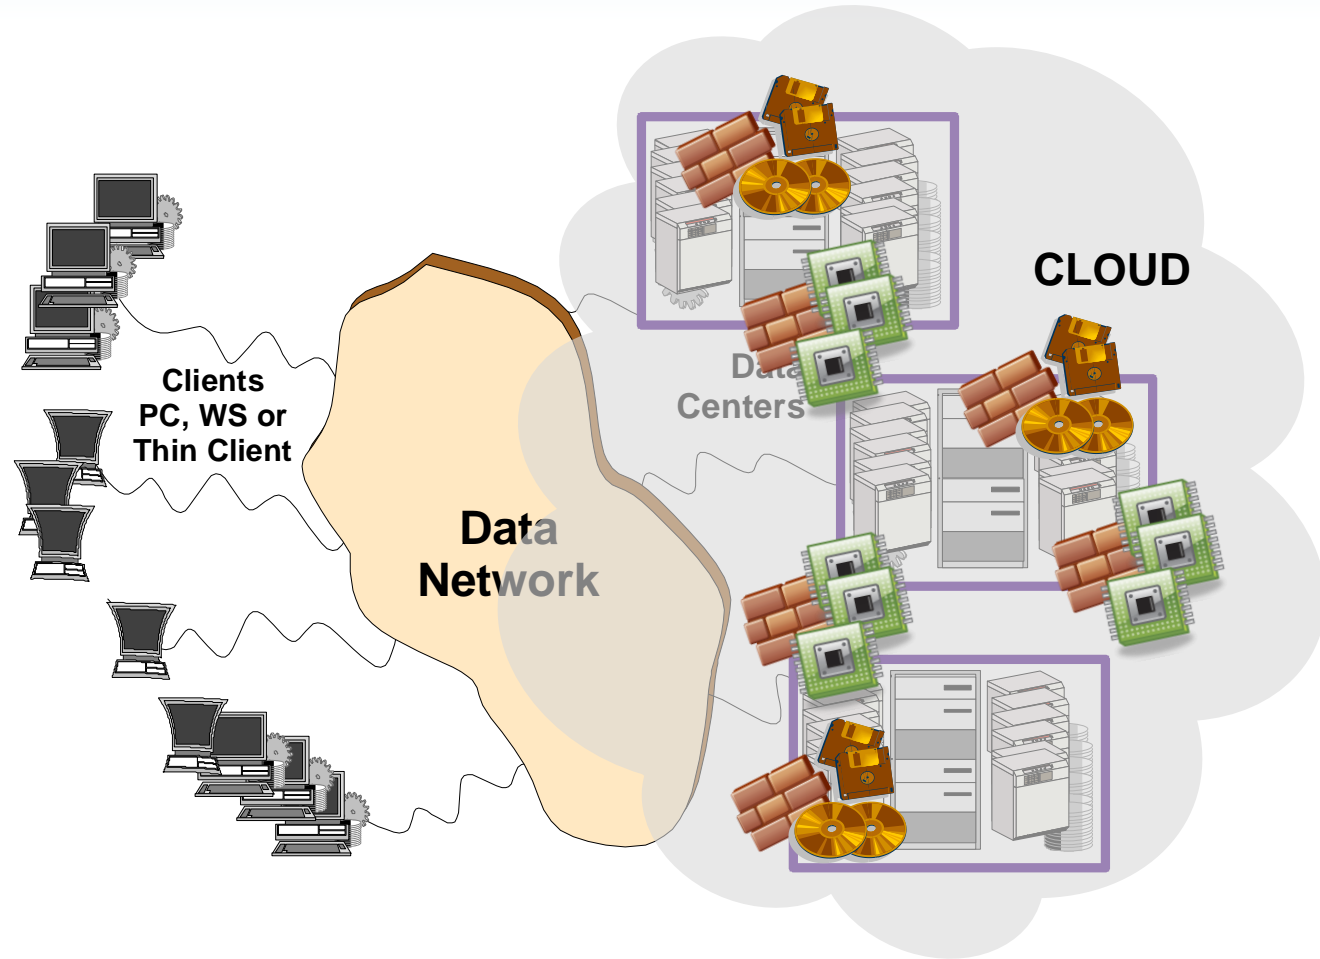


Ciências  
ULisboa





# Cloud Computing



# Rudiments of: Dependability, Real-Time, Security

# Non-functional properties of Distributed Systems

- What systems **are**, rather than what **they do**
- Some familiar names:
  - Dependability and fault-tolerance: **reliability** and **availability**
  - Real-time: **timeliness** and **predictability**
  - Security: **confidentiality** and **integrity**

▪ Dependability	(section 6.1)
▪ Timeliness	(section 11.1)
▪ Security	(section 16.1)



**Ciências**  
**ULisboa**

# Rudiments of: Dependability

# How reliable are distributed systems?



Ciências  
ULisboa

***A distributed system is one that prevents you from working due to the failure of a machine we even had never heard of.***

*[L. Lamport]*

- Since:
  - machines may fail independently (and still...),
  - but influence each other,
  - communicating through non-reliable networks with unpredictable delays
- ... “gathering machines” just worsens situation:
  - reliability ( $<1$ ) of a system is equal to the product of the reliabilities of all individual components needed for the system to work correctly
  - $R(10 \text{ machines} @ 0.99) = 0.99^{10} = 0.90$
  - $R(10 \text{ machines} @ 0.90) = 0.90^{10} = 0.35$

# How do distributed systems fail?

*Why do computers fail and what can we do about it?*

[J. Gray]

- Why:
  - because all that works fails
  - maybe they look like failing too much...
  - but we have a tendency of overestimating our HW and SW
- So?:
  - better prevent than remedy
- Dependability is ...
  - the trustworthiness of a computing system in fulfilling a set of properties, which allows justifiable trust on the service it delivers
- How?
  - by appropriate mechanisms that allow the prevention or tolerance (as well as quantification or prediction) of disturbances

# Threats to dependability

They upset the service execution during the lifetime of the system.

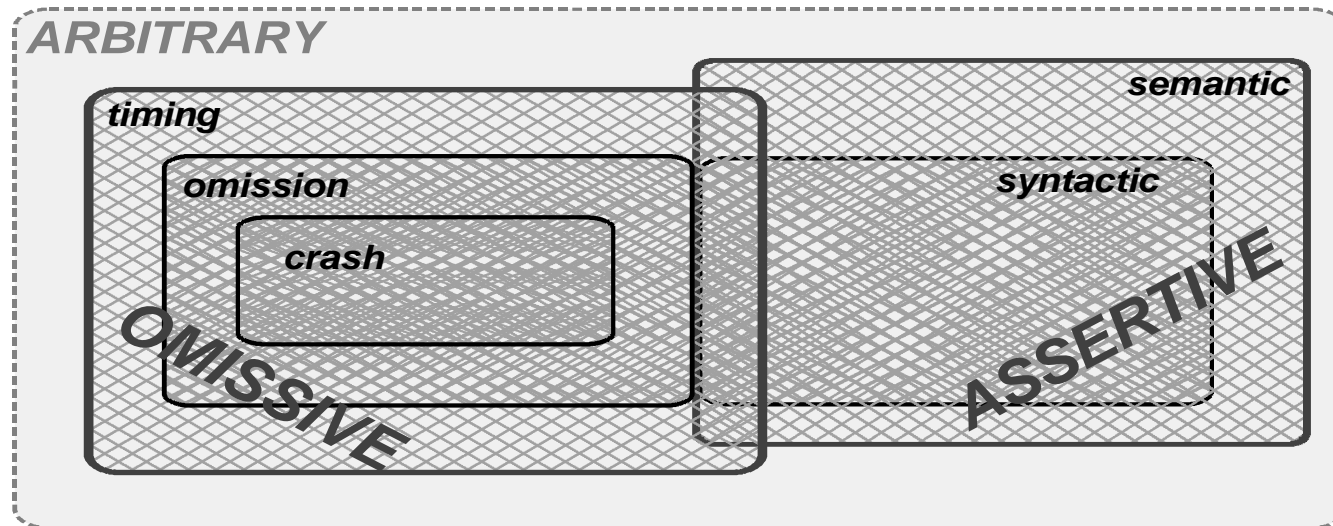
- Types:
  - **Falta (fault)** – cause of an error
  - **Erro (error)** – manifestation of a fault in the state of the system
  - **Falha (failure)** – consequence of an error on the service of system
- Simple Examples:
  - **Fault-** bit stuck at zero in a RAM; a packet is lost in transmission
  - **Error** – reading zero, after having written one; message never arrives
  - **Failure-** the wrongly read bit is delivered to the user; message never delivered to user
- Possible solutions to the examples above:
  - Detection- memory with parity bit; timeout set for packet arrival
  - Recovery- memory with ECC; repeat transmission if timeout expires
  - Masking- replicate memory; systematically repeat a transmission

# Interaction Faults classification

(specially important in distributed systems)



Ciências  
ULisboa



- **Omissive**

- **Crash**
  - Interactions stop
- **Omission**
  - some interactions are missed
- **Timing**
  - interactions out of time (early or late)

- **Assertive**

- **Syntactic**
  - temperature sensor that indicates +ab degrees
- **Semantic**
  - temperature sensor that indicates 26° when it is 30°



# Dependability properties

- **Reliability**

- The measure of the continuous delivery of correct service
- Can be expressed in terms of Mean Time To Failure (MTTF) or **Mean Time Between Failures (MTBF)**

- **Maintainability**

- The measure of the time to restoration of correct service
- Can be expressed in terms of **Mean Time To Repair (MTTR)**

- **Availability**

- The measure of delivery of correct service with respect to alternation between correct and incorrect service
- Can be expressed in terms of a ratio between uptime and total time, or using the previous two metrics: **Availability = MTBF / (MTBF+MTTR)**

- **Safety**

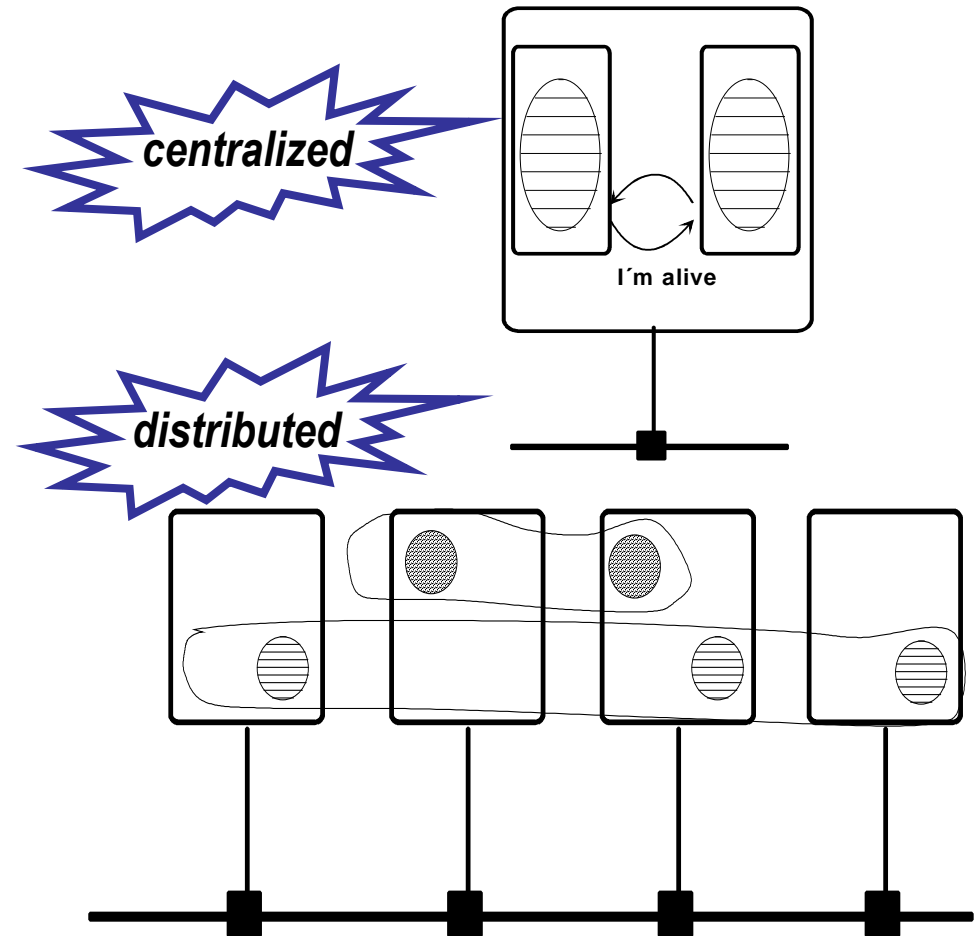
- The degree to which a system, upon failing, does so in a non-catastrophic manner

# Means to Attain Dependability

- Fault **prevention** – how to prevent the occurrence of faults?
  - Avoid design defects (good process; tools; etc.)
  - Shielding, robust design, good operational procedures etc. for runtime faults
- Fault **tolerance** – how to deliver correct service in the presence of faults?
  - Error detection
  - Recovery
  - Fault handling
- Fault **removal** – how to reduce the number/severity of faults?
  - Verification & validation (design phase)
  - Corrective & preventive maintenance (operations)
- Fault **forecasting** – how to estimate the future incidence of faults?
  - Simulation, modeling, prediction from process metrics (data from models)
  - Historical data, accelerated life testing, etc. (data from real systems)

# Foundations of modular and distributed fault tolerance

- **Topological separation**
  - Failure independence
  - Graceful degradation
- **Replication**
  - Software vs. hardware
  - Fine granularity
  - Resource optimization
- **Incremental F/T by:**
  - Class (omissive, semantic)
  - Number of faults
  - Number of replicas
    - pairs, triples, etc.
  - Type of replica control
    - active, passive
    - round robin, voting





**Ciências**  
**ULisboa**

# Rudiments of: Real-Time

# Real-Time (RT) Systems

- **Real-time (RT) system:** system which has at least one service whose progression is specified in terms of timeliness requirements dictated by the environment
  - The “real-time problem” is dictated by the need to synchronize our actions with the environment.
  - Real-time is about predictability (timeliness)
- Some **misconceptions:**
  - RT is ad-hoc design, assembly program., interrupts, and so forth;
  - RT systems are automata, pre-programmed and static;
  - RT is about having enough speed, and ever-increasing MIPS and Mbauds will solve all “performance” problems;
  - RT deadlines do not make sense, since they will be missed because failures occur, messages get lost, software has bugs, etc.
- **Examples of real-time systems:**
  - Oven controller, Manufacturing cell, Traffic Light Controller, Air traffic control system

# Classes of RT Systems

## importance of scheduling



Ciências  
ULisboa

- Hard real-time systems, where timing failures are to be **avoided**
  - E.g. on-board flight control system (fly-by-wire)
- Soft real-time systems, where **occasional** timing failures are accepted
  - E.g. on-line flight reservation system
- Mission-critical real-time systems, where timing failures **should be avoided** and occasional failures are handled as exceptional events
  - E.g. air-traffic control system

# Illustrating the scheduling problem: the Fable of the Hare and the Turtle...

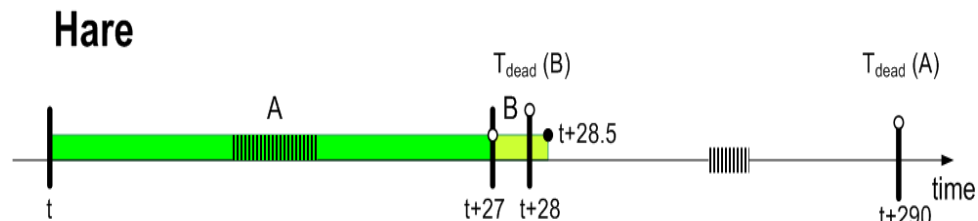


Ciências  
ULisboa

*Fable of La Fontaine, adapted to a R/T scheduling problem by G. LeLann*

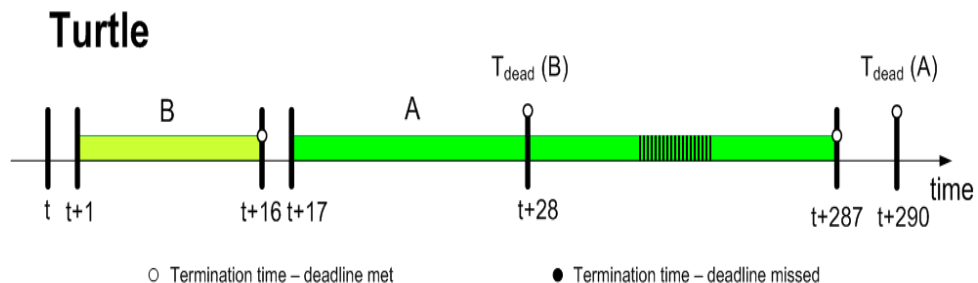
- HARE System:

- Speed = 10;
- Context switch = 0
- Scheduling: Task A is started first



- TURTLE System:

- Speed = 1;
- Context switch = 1
- Scheduling: Task B is started first



- Problem Definition:

- Task A: duration= 270; deadline= t+290
- Task B: duration= 15; deadline= t+28
- Tasks A and B released at same time

# Illustrating the scheduling problem: the Fable of the Hare and the Turtle...

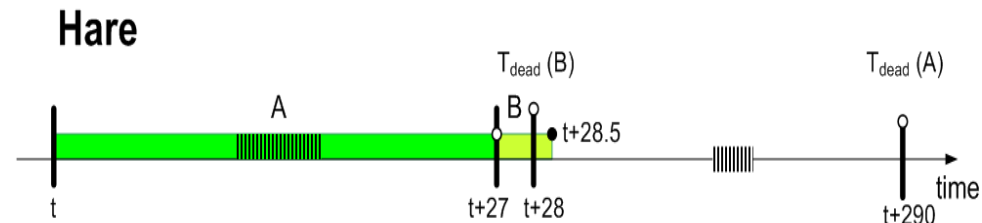


Ciências  
ULisboa

*Fable of La Fontaine, adapted to a R/T scheduling problem by G. LeLann*

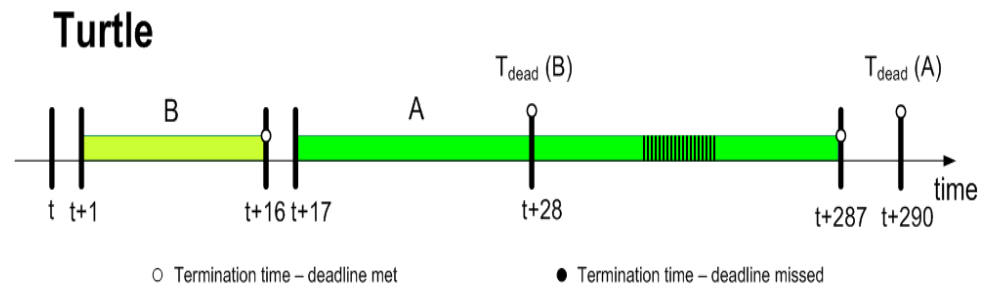
- HARE System:

- Speed = 10;
- Context switch = 0
- Scheduling: Task A is started first



- TURTLE System:

- Speed = 1;
- Context switch = 1
- Scheduling: Task B is started first



○ Termination time – deadline met

● Termination time – deadline missed

- Problem Definition:

- Results:

- Task A: duration= 270; deadline= t+290
- Task B: duration= 15; deadline= t+28
- Tasks A and B released at same time
- **HARE, optimistic, starts to serve A and misses the deadline of B**
- **TURTLE, realistic, starts to serve B (closer deadline) and meets both deadlines.**





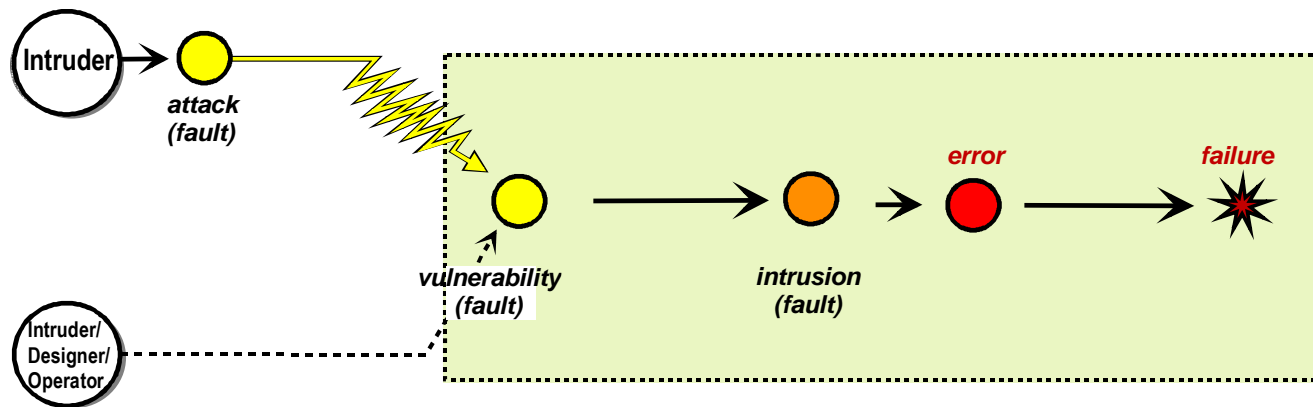
**Ciências**  
**ULisboa**

# Rudiments of: Security

# Insecurity, People and Computers

- Insecurity is as concerned with technical deficiencies as with **people's attitudes**
- Insecurity is quantitatively caused to a great extent by the actions of people who must be **educated** about the seriousness of their deeds
- It is better to invest than to spend: **improve your system before intrusion**
- **Don't use crypto "just because"**: cryptocracy (a form of technocracy) is enemy of good systems practice

# A systematic view of security failure sequence Attack-Vulnerability-Intrusion



AVI sequence : *attack + vulnerability → intrusion → error → failure*

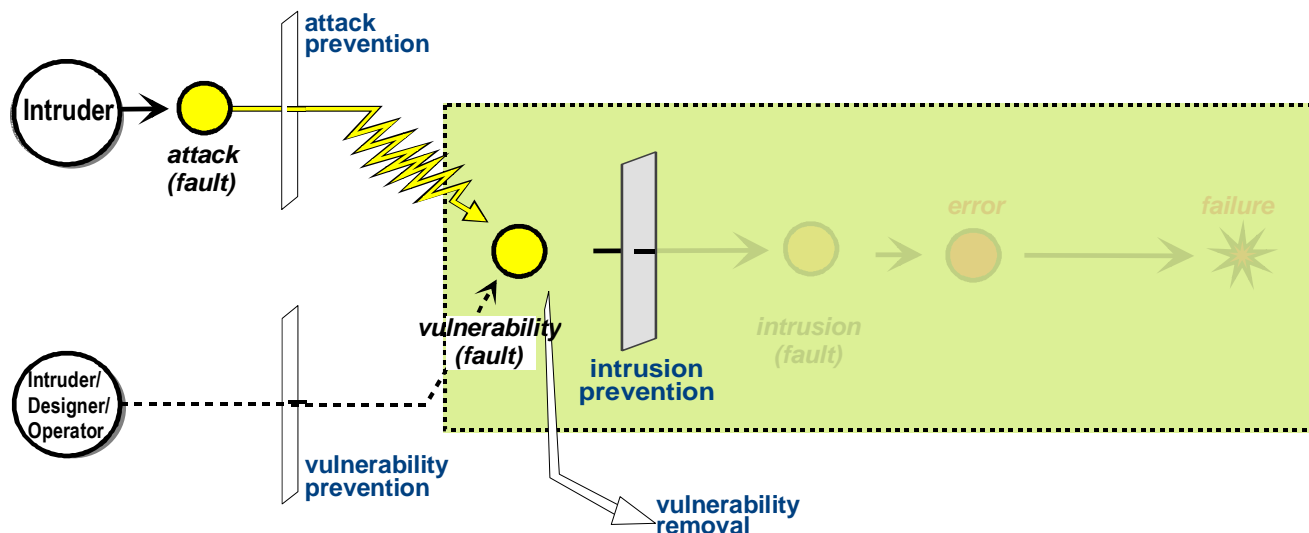
# Security Properties

- **Confidentiality**
  - The measure in which a service or piece of information is protected from unauthorized disclosure
- **Integrity**
  - The measure in which a service or piece of information is protected from illegitimate and/or undetected modification
- **Authenticity (\*)**
  - The measure in which a service or piece of information is genuine, and thus protected from personification or forgery
- **Availability**
  - The measure in which a service or piece of information is protected from denial of authorized provision or access

(\*) Considered integrity of meta-information by some authors

# Avoiding security failure and achieving security properties

## canonical track: intrusion prevention



sequence : *attack + vulnerability* → *intrusion* → *failure*



**Ciências**  
**ULisboa**

# Some formal notions

# Formal system predicates

## a colloquial definition

- **Objective:**
  - Specify in a formal manner, including formulas containing logic (*and, or, exists, forall*), temporal logic (*eventually, always*) and time (*until/from*) operators
- We can specify the properties of any program or protocol in terms of properties of: **safety** and/or **timeliness** , **liveness**

# Formal system predicates

## a colloquial definition

- **Safety**
  - the measure in which a service/program **does not do bad things**
  - safety properties specify that wrong events never take place
  - a safety property specifies that a predicate  $P$  is always true
  - example, "any delivered message is delivered to all correct participants" is a safety property
  - If it is not secured, the system becomes incorrect
  - however, it does not impose that messages are delivered at all...
  - So, what do you think is missing here?



# Formal system predicates

## a colloquial definition



Ciências  
ULisboa

- **Liveness**

- the measure in which a service/program **does good things**
- liveness properties specify that good events eventually take place
- a liveness property specifies that predicate  $P$  will eventually be true
- example, "any message sent is delivered to at least one participant" is a liveness property
- If it is not secured, the system may not progress (messages are not delivered)
- With the previous safety property, "any delivered message is delivered to all correct participants" , the system might not do anything at all and still be correct!
- Together with the current liveness property, "any message sent is delivered to at least one participant" , the system will both be correct and make progress!
- Why couldn't we have said: "any message sent is delivered to all correct participants" ??

# Formal system predicates

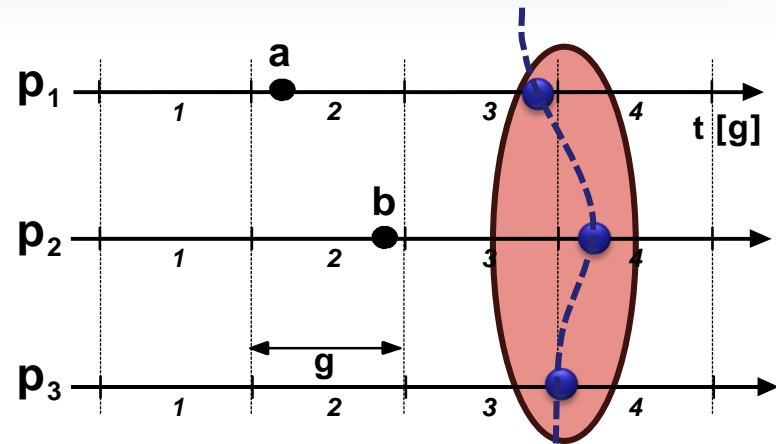
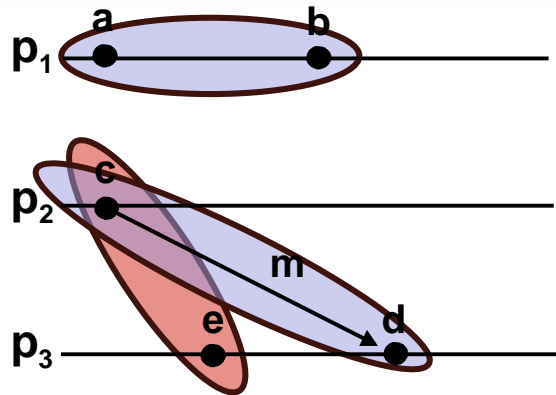
## a colloquial definition



Ciências  
ULisboa

- **Timeliness**
  - a sub-class of safety property is timeliness, which specifies that **a predicate  $P$  will be true in relation to a given instant of real time** (until, before, at)
  - "any transaction completes until  $T_t$  from the start" is a timeliness property
  - for the property to be secured, all transactions must execute within  $T_t$  time units

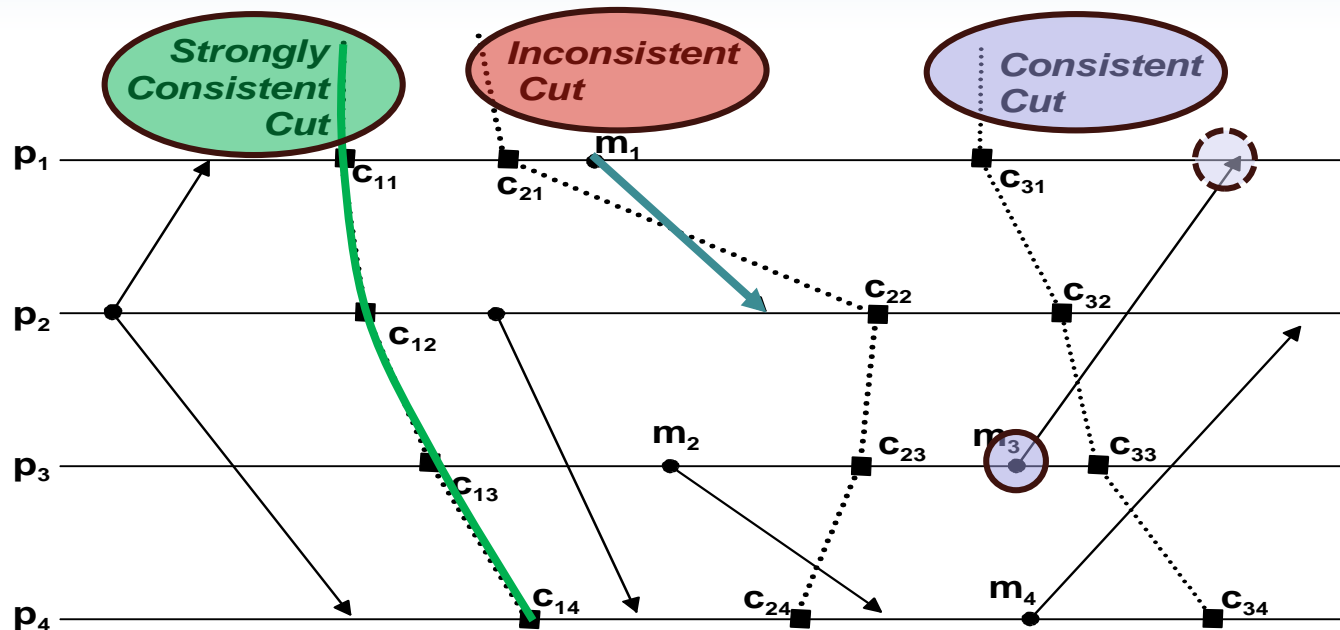
# Space-Time and Lattice Diagrams



real clocks  
differ

- Event types:
  - local execution, send, receive, deliver
- Precedence:
  - $a < b ; c < d$
  - is  $c < e$  ?
- Time lattices:
  - based on notion of global time with granularity (tick)  $g$
  - does 4:00 tick at exactly the same place everywhere?

# Cuts and Global States (GS)



- Types of cuts:
  - inconsistent** cut: snapshot gives invalid picture of GS
  - consistent** cut: snapshot gives correct but possibly incomplete picture of GS (e.g., ignores messages in transit)
  - strongly consistent** cut: snapshot faithfully represents GS