



Ciências
ULisboa

Programação em Sistemas Distribuídos

MEI-MI-MSI

2018/19

3. Models of Distributed Computing

Prof. António Casimiro

Distributed Computing Models

Main models, neither exhaustive nor air-tight




Ciências
ULisboa

- Client-Server (RPC, RMI, WWW) (Cliente-Servidor)
- Distributed Objects (Objectos Distribuídos)
- **Distributed Shared Memory (DSM, Tuples) (Memória Partilhada Distribuída)**
- Distributed Atomic Transactions (Transacções Atómicas Distribuídas)
- Message-oriented (Message Queue, Publish/Subscribe) (Orientado para mensagens, Fila de Mensagens, Editor/Assinante)
- Stream (Corrente)
- Group-Oriented (Orientado para Grupos)
- Peer-to-peer (Inter-pares)

Distributed Shared Memory (DSM)

Virtues of DSM

- Allowing the shared memory model to be used in systems that don't have physical shared memory
- Allowing all nodes to share a virtual address space (through the shared memory model)
- Allowing communication costs to be reduced by maintaining data copies on the location of access




Desirable characteristics of DSM

- Transparency
 - Memory distribution should be transparent to the application and give the illusion of a centralized shared memory system (semantics and performance)
- Traffic optimization
 - Algorithms should minimize number of messages exchanged
- Remote access masking
 - Remote accesses should be hidden from programmer (e.g. node forced to wait for a message from another node in order to proceed)
- False sharing prevention
 - Minimize probability of contention of nodes caused by page granularity (e.g., when two nodes update unrelated variables that, by coincidence, are located in the same page)

DSM Architectures

- **Centralized server**

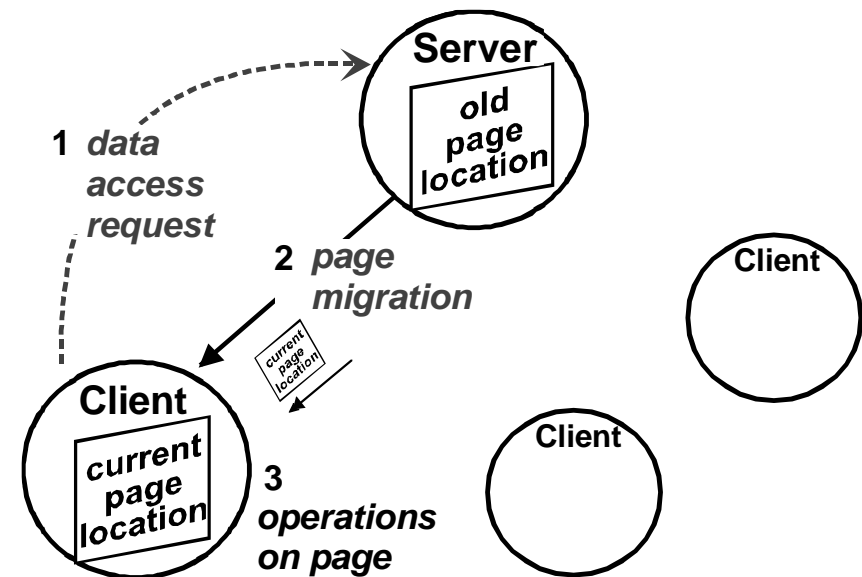
- Naive architecture, for starters
- Central server holds the memory pages and all data accesses are performed invoking the server
- Server serializes all requests, thus ensuring strong consistency
- System behaves just like if a single central memory was available
- Correctness relies on having a single central memory on the server!
- **Cons:** since memory accesses are remote, performance is deplorable



DSM Architectures

- **Page migration**

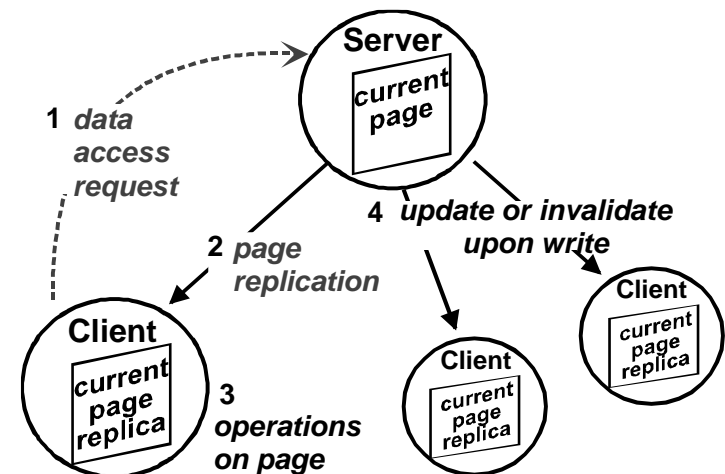
- One way to improve on central server architecture is exploring locality
- Instead of forcing all accesses (reads, writes) to execute remotely, migrate a chunk of memory (page, block) from the server to the client
- Only one node can access the shared data chunk at a time
- Subsequent accesses to that same page performed locally by the client
- If another client later wants to access the same page, the page is migrated again
- Correctness depends of a given page being at a single location at any given point in time
- **Cons:** high client contention for a page causes it to go back and forth (trashing), performance is affected



DSM Architectures

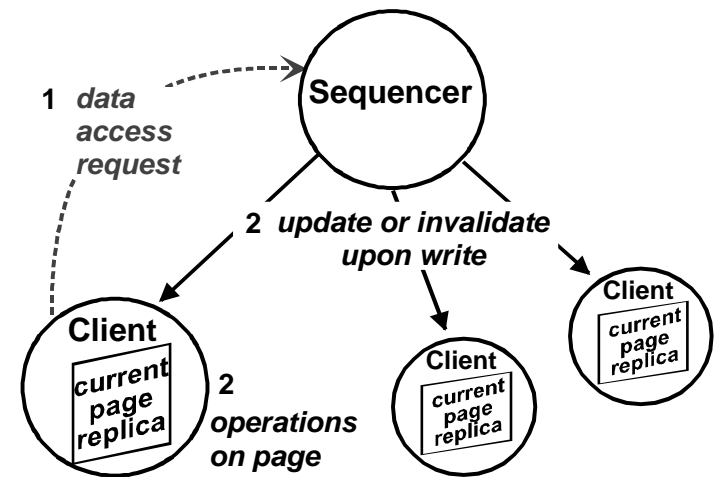
- **Read replication**

- Assume reads are much more common than writes
- **Read-only replicas** held in several clients
- **Read-write replica** held in a single clients
- After a write, all copies are invalidated or updated
- **Write-invalidate**: all shared copies are invalidated before write executes
- **Write-update**: all shared copied are updated with new value




DSM Architectures

- **Full replication**
 - Assume reads are much more common than writes
 - **Read-write permanent replicas** held in several clients
 - Write requests are coordinated by a sequencer: clients perform competitive access requests which are ordered



DSM Architectures

- **Full replication with distributed sequencer**
 - Sequencer may be decentralized, implemented by a distributed algorithm
 - Most elegant and widely used approach



DSM consistency

Fundamental models



Ciências
ULisboa

- **Strict consistency or atomic consistency:**
 - P1- Every read to a data x returns the value most recently written to x
 - P2- Writes and reads are ordered according to the **real time** order
 - In practice not fully achievable, due to clock uncertainties
- **Sequential consistency:**
 - P1- Every read to a data x returns the value most recently written to x
 - P2- Writes and reads are ordered according to the **program** order
 - Weaker than strict consistency
 - Achievable with total ordering of writes
- **Causal consistency:**
 - Weaker than sequential consistency
 - Only write operations that are causally related must be ordered
 - Therefore, no total ordering is needed for concurrent writes

DSM consistency

Fundamental models



Ciências
ULisboa

- **FIFO consistency:**
 - Weaker than causal consistency
 - Operations must only be ordered respecting process order
 - Example (assume $x=0$ in the beginning):
 - P1 writes $x=1$
 - P2 writes $x=2$
 - P3 can read, in sequence, $x=1$ and then $x=2$
 - P4 can read, in sequence, $x=2$ and then $x=1$
 - Would this also be a valid sequence if the DSM offers causal consistency?
 - What about if P2 reads $x=1$ and then writes $x=2$?
- **Release consistency:**
 - Explicit synchronization with acquire (i.e. lock) and release (i.e. unlock) operations
 - Programs lose transparency, but no DSM mechanisms need to be applied while applications are within critical sections (memory may become inconsistent in these periods)
 - Consistency enforced upon release operations



Ciências
ULisboa

Tuple Spaces

Characteristics of Tuple Spaces

- Processes communicate indirectly by placing tuples in a tuple space, which other processes can read or remove
- Tuples do not have an address but are accessed by pattern matching on content (content-addressable memory)
- Industrial implementations: Sun JavaSpaces, IBM Tspaces

Tuple space

Generic Interface



Ciências
ULisboa

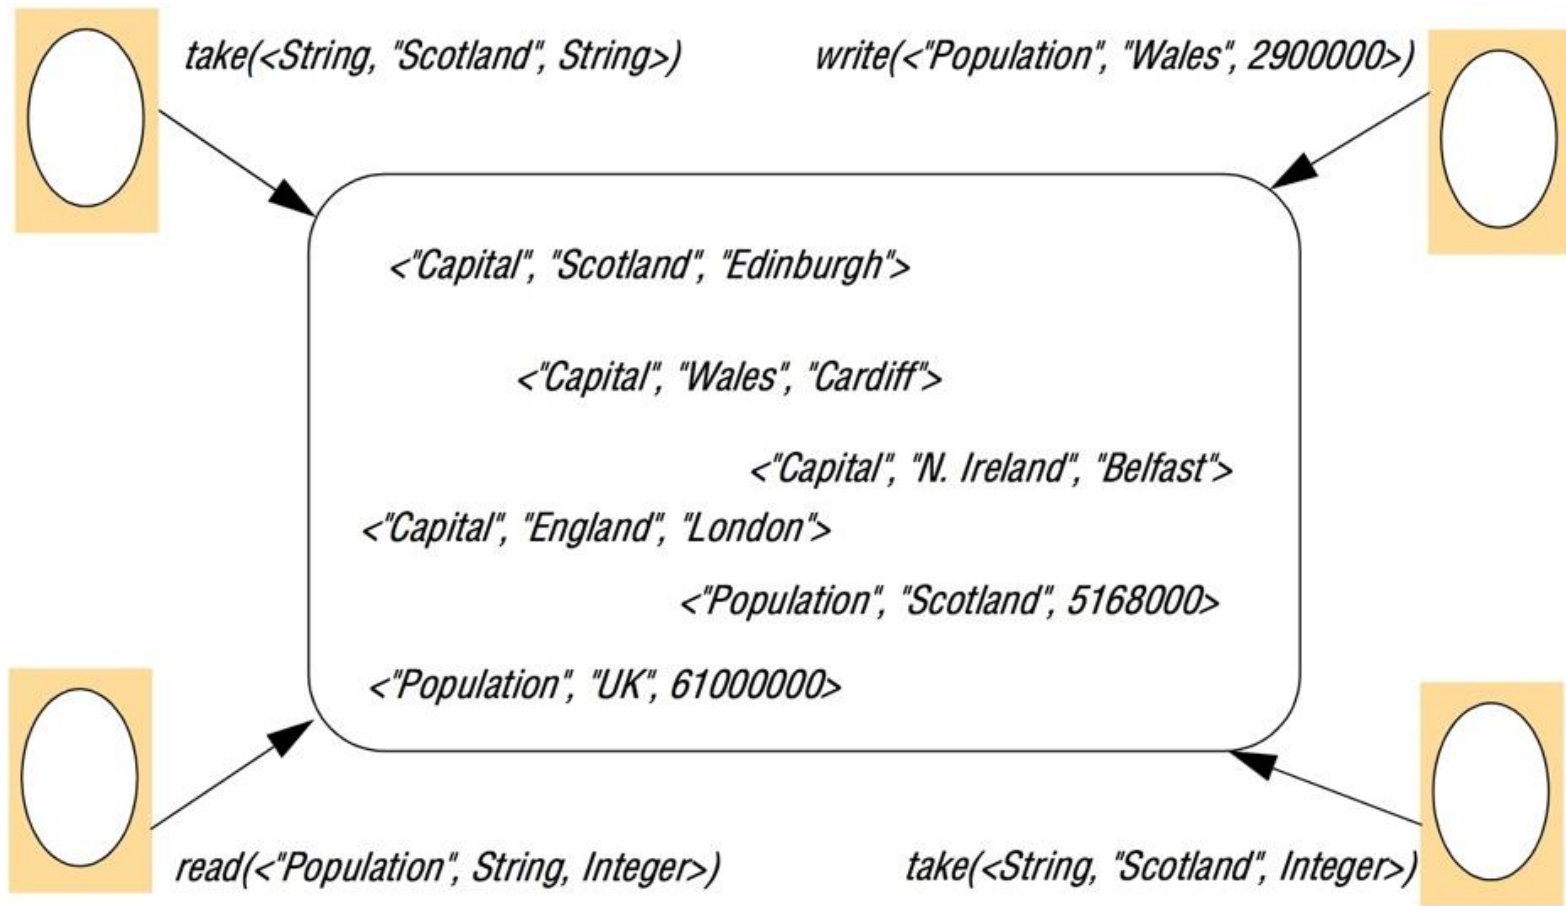
- ***write(t)***
 - Requesting process inserts the tuple ***t*** into the tuple space
- ***read(tspect)***
 - Requesting process issues read request to tuple space
 - Tuple(s) matching requested tuple specification ***tspect*** are returned, space is not affected
- ***take(tspect)***
 - Requesting process issues take request to tuple space
 - Tuple(s) matching requested tuple specification ***tspect*** are returned and extracted from space

Tuple space

Programming model



Ciências
ULisboa



Tuple space

Programming model



Ciências
ULisboa

- Tuple specification
 - Read or take provide a tuple specification `tspec` and the tuple space returns any tuple that matches that specification
- Blocking
 - Read and take operations both block until there is a matching tuple in the tuple space
- Tuples are immutable
 - No direct access to tuples in tuple space is allowed
 - To modify tuples, processes replace them: take, modify, then write
- Programming examples:
 - `take(<String, "Scotland", String>)` will match `<"Capital", "Scotland", "Edinburgh">`
 - `take(<String, "Scotland", Integer>)` will match `<"Population", "Scotland", 5168000>`
 - `write(<"Population", "Wales, 29000000>)` inserts a new tuple in the tuple space with information on the population of Wales