



Ciências
ULisboa

Programação em Sistemas Distribuídos

MEI-MI-MSI

2018/19

3. Models of Distributed Computing

Prof. António Casimiro

Distributed Computing Models

Main models, neither exhaustive nor air-tight



Ciências
ULisboa

- Client-Server (RPC, RMI, WWW) (Cliente-Servidor)
- Distributed Objects (Objectos Distribuídos)
- Distributed Shared Memory (DSM, Tuples) (Memória Partilhada Distribuída)
- Distributed Atomic Transactions (Transacções Atómicas Distribuídas)
- Message-oriented (Message Queue, Publish/Subscribe) (Orientado para mensagens, Fila de Mensagens, Editor/Assinante)
- **Stream (Corrente)**
- Group-Oriented (Orientado para Grupos)
- Peer-to-peer (Inter-pares)



Ciências
ULisboa

Stream

Stream models

Applications



Ciências
ULisboa

- **Web-based multimedia:**
 - Best-effort quality of service for access to streams of audio and video data published via Web
 - Multimedia data retrieved from large online storage systems
 - Applications such as YouTube, Netflix and BBC iPlayer
- **Audio/Video conferencing:**
 - Full-duplex end-to-end audio/video conferencing
 - Applications such as Skype (Microsoft), Hangouts (Google), iChat (Apple)
- **Live streaming:**
 - Requires dedicated network bandwidth for good quality and dedicated video server
 - Dedicated or web-based clients

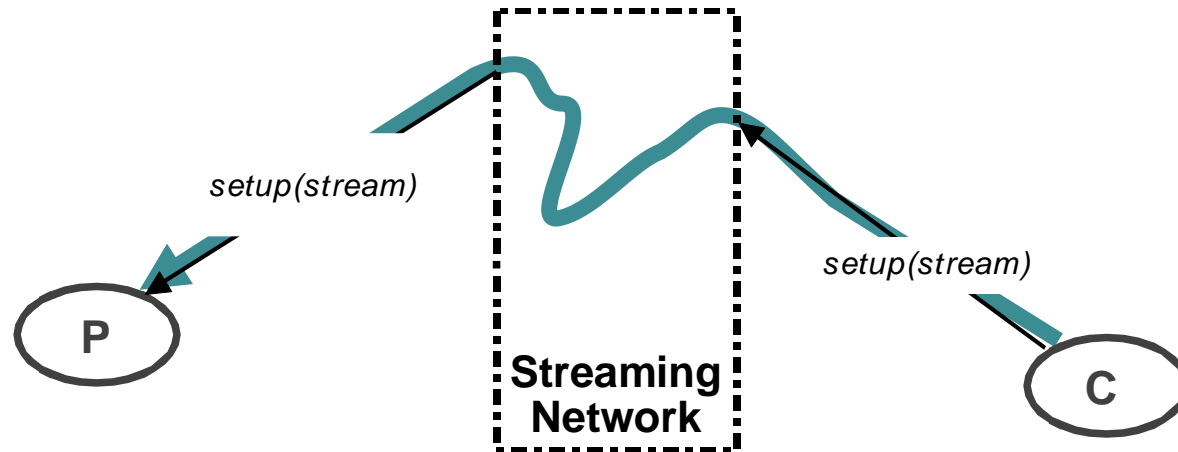
Stream models

- Can be seen as **special case of channel-based publish/subscribe**
- Space decoupling
 - Source (producer) and sink (consumer) do not have to know each other
- Time coupling
 - Producer and consumer **must be simultaneously active**
 - Stream is a flow of information obeying **real-time constraints** from the end-user viewpoint
 - In content distribution networks (CDN), producer holds repository
- Synchronization coupling
 - Producers and consumers are synchronized to enforce the stream abstraction
- Stream delivery:
 - Stream is divided in ordered chunks, which are transmitted through the network
 - R/T message dissemination protocols enforcing message delivery deadlines make sure that each chunk arrives in time to be collated at the right place in the timeline
 - Client side buffering helps mitigating network latency jitter

Streaming Information flow




Ciências
ULisboa



Streaming Information flow




Ciências
ULisboa



Streaming Network

Architectural alternatives


- Baseline Internet
 - R/T protocols (RTP)



Streaming Network

Architectural alternatives

- Baseline Internet
 - R/T protocols (RTP)
- R/T Overlay Network
 - streaming brokers



Distributed Computing Models

Main models, neither exhaustive nor air-tight



Ciências
ULisboa

- Client-Server (RPC, RMI, WWW) (Cliente-Servidor)
- Distributed Objects (Objectos Distribuídos)
- Distributed Shared Memory (DSM, Tuples) (Memória Partilhada Distribuída)
- Distributed Atomic Transactions (Transacções Atómicas Distribuídas)
- Message-oriented (Message Queue, Publish/Subscribe) (Orientado para mensagens, Fila de Mensagens, Editor/Assinante)
- Stream (Corrente)
- **Group-Oriented (Orientado para Grupos)**
- Peer-to-peer (Inter-pares)



Ciências
ULisboa

Group-oriented

Group-based communication

Recall:

- RPC is point-to-point and asymmetric
 - How to handle multi-participant interactions?
 - E.g. fragmented and/or replicated database
 - How to handle (symmetric) conversations?
 - E.g. email-group, video conversations, collaborative on-line file editing
- RPC is a remote operation, in closed-circuit, blocking, for reliability
 - How to achieve reliability in open-circuit?
 - E.g. with non-blocking but reliable operations
- **Group-based communication has interesting properties to solve the problems mentioned above**

Group-oriented systems

Characteristics



Ciências
ULisboa

- **Group**: set of entities that may receive messages
 - E.g. group of machines, group of processes
- **Group address**: the group is addressed collectively
 - E.g. list of addresses, multicast address, logical address
- **Group membership**: management of joining and leaving actions and information about membership
 - E.g. who is in the group (view), synchronization of join requests, failure detection (forced leave)
- **Reliable group communication**: usually with high QoS semantics
 - E.g. reliable delivery; ordered delivery; causal delivery

Group-Oriented Systems

Architecture - site vs. participant modules



Ciências
ULisboa

- **Activity Support**

- Provides support for distributed paradigms (e.g. key or replication management)

- **Participant Failure Detector**

- Marks as unreachable the failed participants
- Or participants at a site detected unreachable by Site Membership Service

- **Participant Membership**

- Provides membership information at the participant level

- **Site Membership Service**

- Information about sites participating in group communication

- **Group Communication**

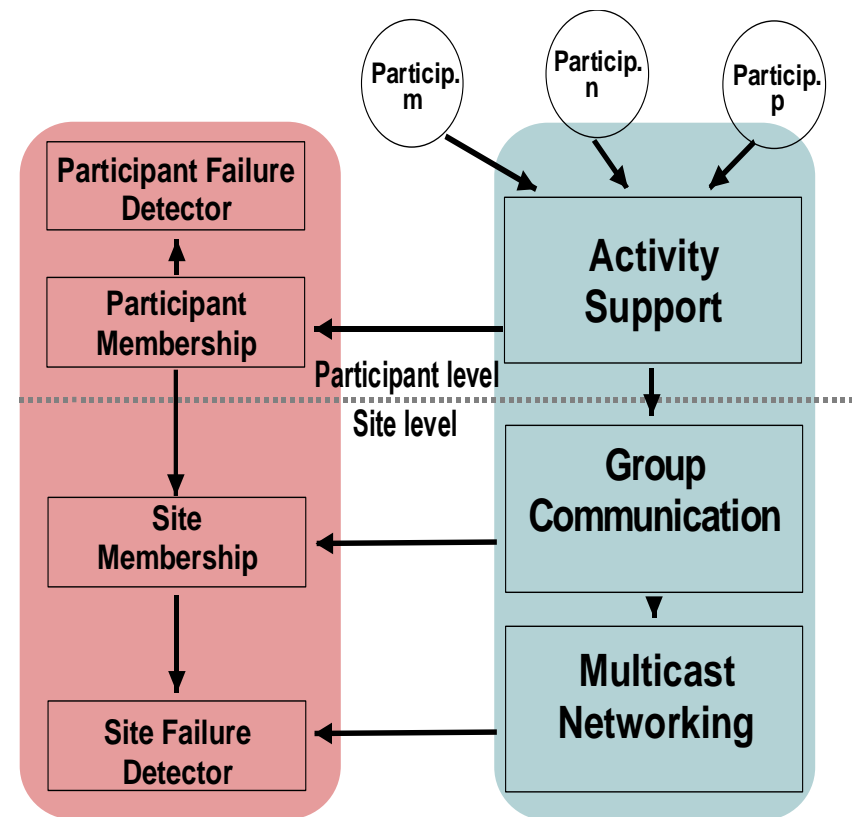
- Provides reliability and ordering guarantees to messages exchanged among sites
- Uses Site Membership to obtain the list of active sites

- **Site Failure Detector**

- Detects failure of other sites

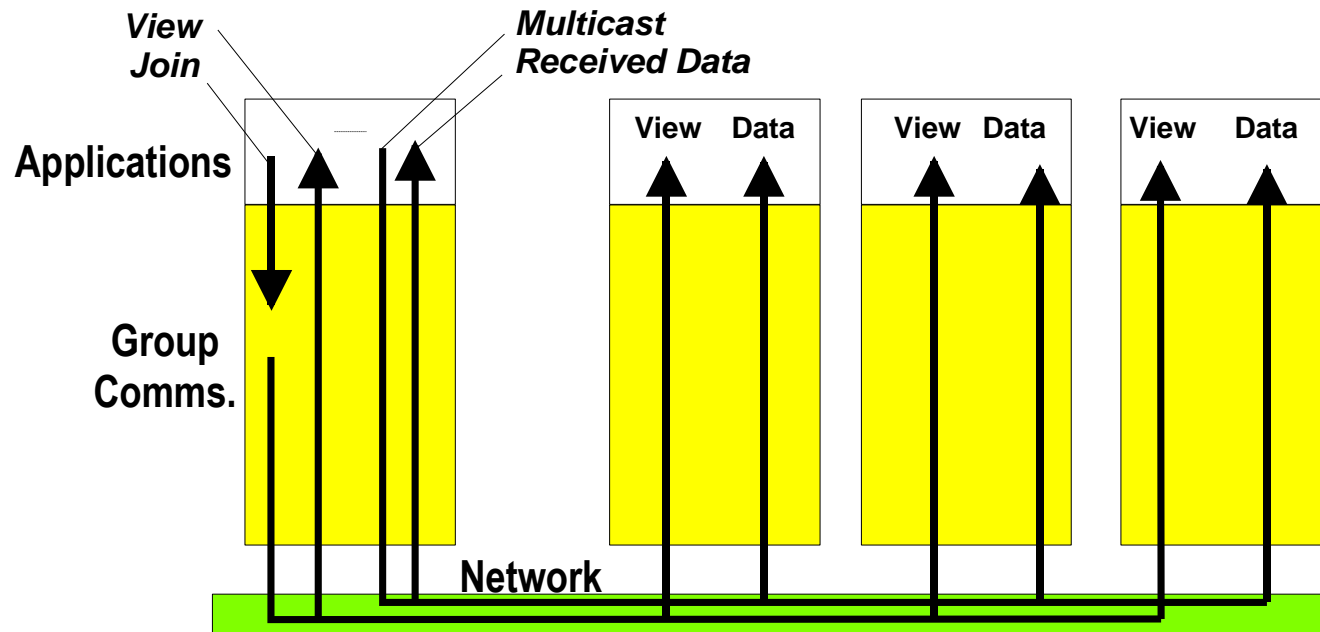
- **Multicast Networking**

- Supports unreliable message multicast services
- Uses Site failure detection to update addressing



Groups in action

- Participant becomes member of a group in response to **join** request
- Participant receives the membership and **view** of the group
- Membership information is also **updated** at all the other members
- As member, participant can **send** to and **receive data** from the group

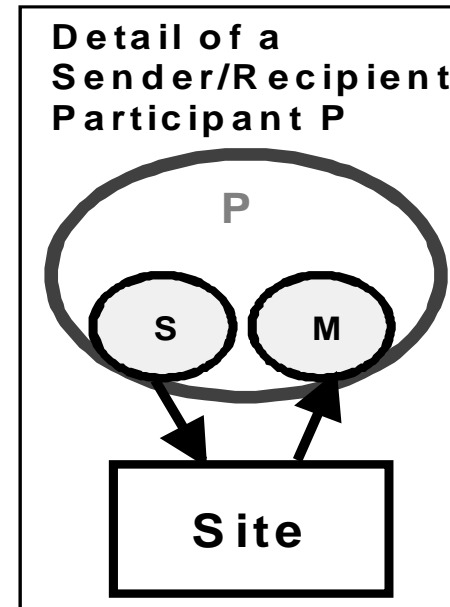



Modelling group access

Unidirectional interaction

- Unidirectional communication
 - Open-circuit
- Full-duplex group access
 - Circuit closed in the process

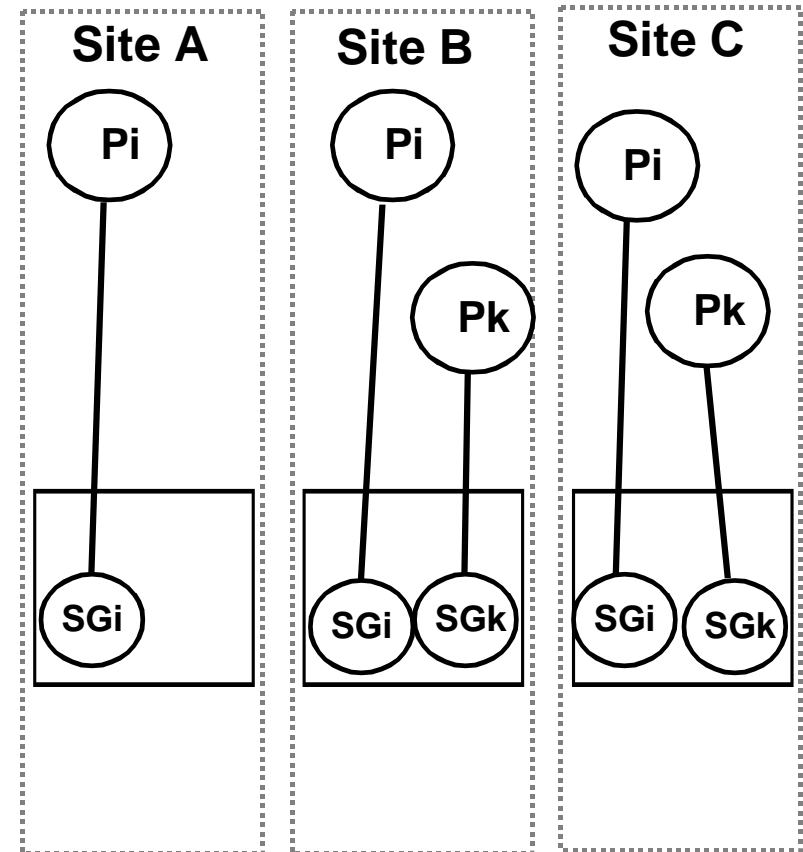
Useful to distinguish *sender* and *member*; *members* receive messages



Group access methods

Normal

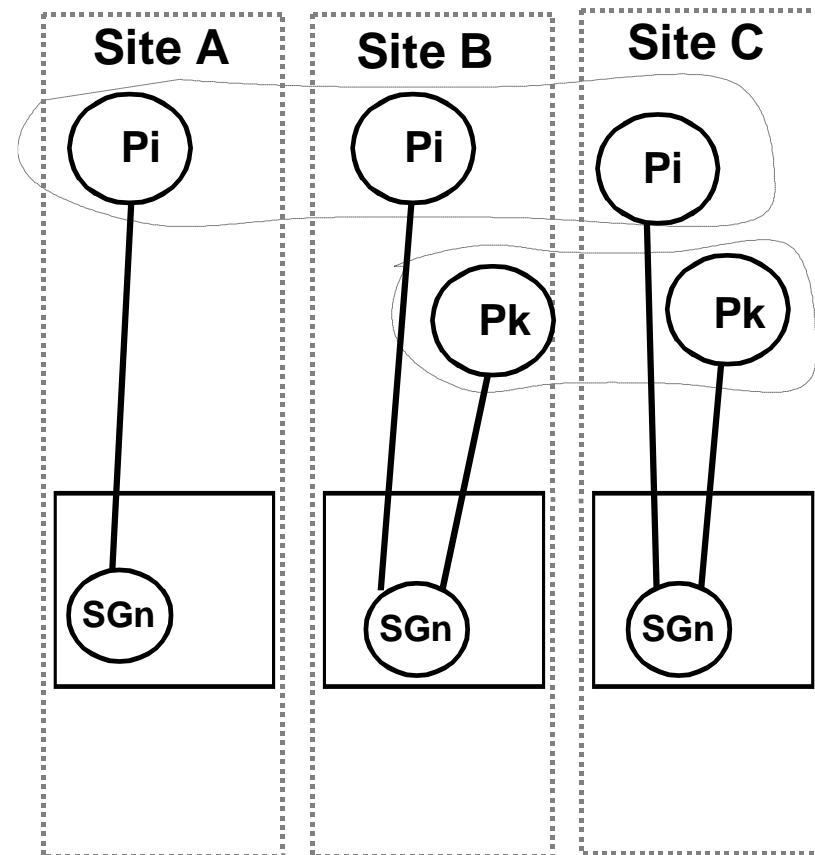
- Participant-level groups P_i and P_k , each supported by site-level groups SG_i and SG_k



Group access methods

Lightweight

- Lightweight groups do resource sharing when several application-level groups have similar membership
- P_i and P_k are now supported by same site-group SG_n



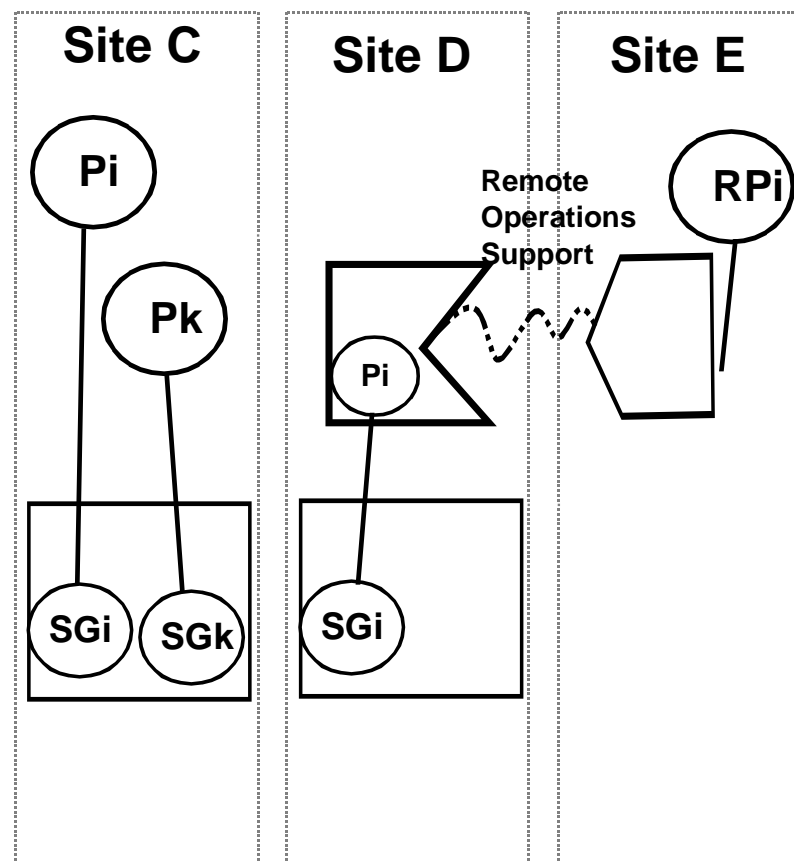
Group access methods

Remote



Ciências
ULisboa



- Senders can be:
 - Groups members
 - Processes outside a group
- **Closed group**
 - Only allows group members to send to the group
- **Open group**
 - Allows non-members to send to the group, called detached senders
- Detached senders must interact with some proxy (group contact): a member or a sender to a group
- Remote sender RP_i communicates with group P_i through proxy at Site D



Group management

- Keep track of nodes or participants
- Dynamic group join/leave
- Failure detection
- Application support:
 - Short addresses (e.g. order number within group)
 - Join/leave/failure notification
 - Consistent group view: the same for all participants
 - View synchrony: total order between view changes and sets of delivered messages

View synchrony



Interaction models with groups

- Client-server
- Conversations
- Dissemination


Interaction models with groups

Client-server



Ciências
ULisboa


- Objectives:
 - Fault tolerance
 - Load balancing
- Characteristics:
 - Clients see **one** logical server, referred by the group address
 - Clients don't know location and number of physical servers
 - Consolidation in origin or destination



Interaction models with groups

Conversations


- Objectives:
 - Symmetric and peer-to-peer activities
- Characteristics:
 - Members with the same role
 - All send and all receive
 - Small groups achieve better QoS



Interaction models with groups

Dissemination

- Characteristics:
 - Producer sends messages to group of consumers
 - Producer sees a logical group of consumers
 - Producer does not need to know number and location of consumers (which might grow)
 - E.g. Multimedia transmission over the Internet, Event publishing



Distributed Computing Models

Main models, neither exhaustive nor air-tight




Ciências
ULisboa

- Client-Server (RPC, RMI, WWW) (Cliente-Servidor)
- Distributed Objects (Objectos Distribuídos)
- Distributed Shared Memory (DSM, Tuples) (Memória Partilhada Distribuída)
- Distributed Atomic Transactions (Transacções Atómicas Distribuídas)
- Message-oriented (Message Queue, Publish/Subscribe) (Orientado para mensagens, Fila de Mensagens, Editor/Assinante)
- Stream (Corrente)
- Group-Oriented (Orientado para Grupos)
- **Peer-to-peer (Inter-pares)**

Distributed Peer-to-Peer (P2P)

Main characteristics of P2P systems

- Distributed
- Decentralized
- Self-organizing
- Dynamic
- Resilient
- Anonymous
- Load balanced
- Symmetric roles
- Resource-sharing



Distributed Peer-to-Peer Model

- Computing with an unreliable and heterogeneous large collection of nodes, with unusual properties:
 - Are **symmetric** in role, unlike Client-Server model
 - Contribute themselves resources: bandwidth, content, storage, memory, CPU (like clients can also be servers and even routers)
- No centralized control, **self-organizing**
 - Nodes are autonomous from central servers or administration
 - Communicate and collaborate directly with each other (not through well-known servers or official routes)
 - Dynamic environment, **frequent join and leave (churn)**
- Inherently resilient
 - No central single point of failure
 - Geographic separation
 - Replication for fault-tolerance easily included and adapted
- Nodes may have widely varying capabilities
 - Storing, content delivery, computing, etc.

Canonical P2P application examples

A snapshot

- Napster
- SETI@HOME
- Video Streaming
- Gnutella



Ciências
ULisboa

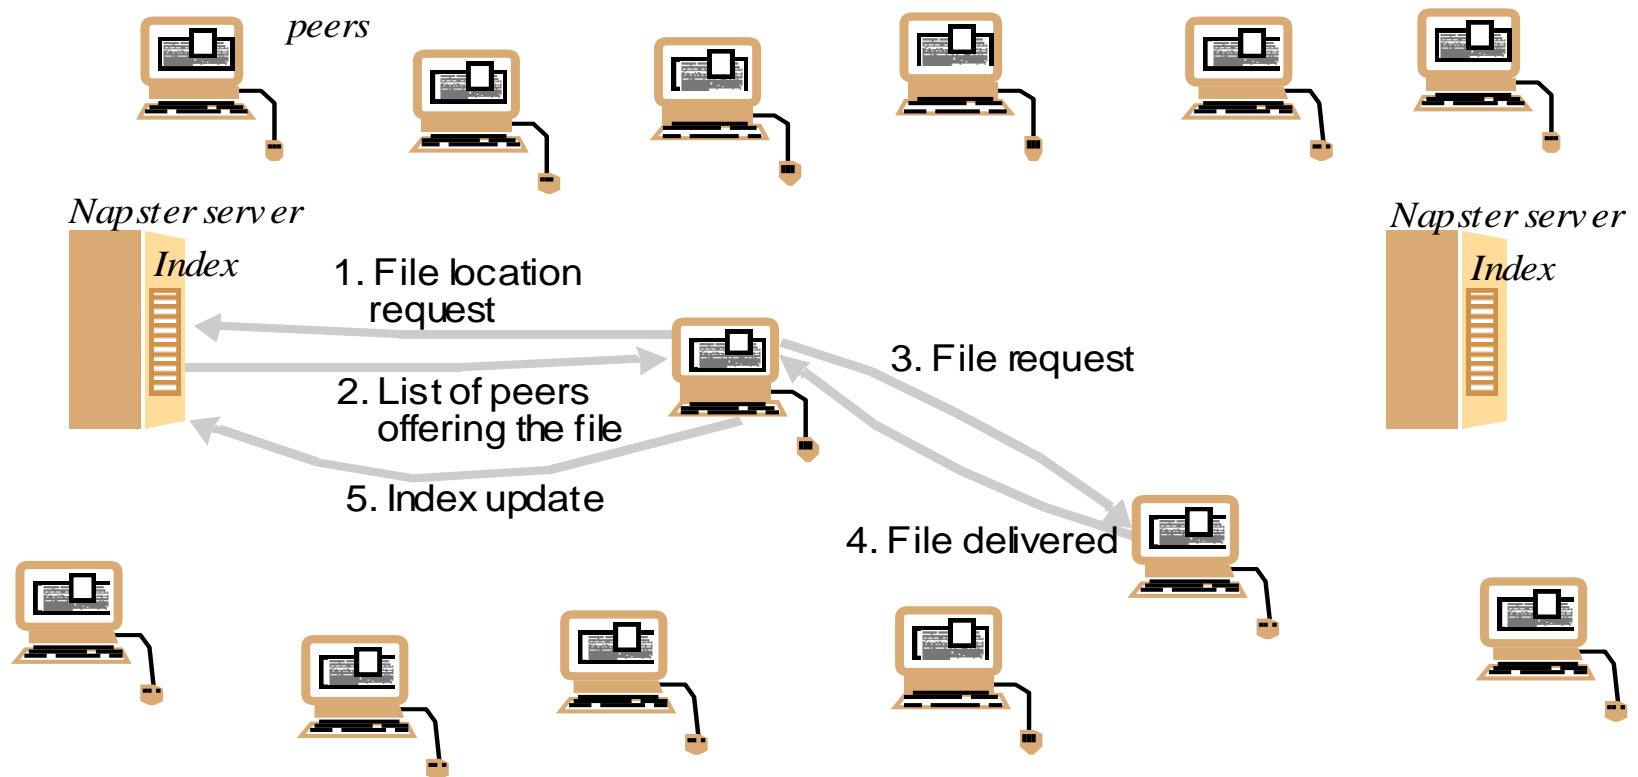
Napster

Peer-to-peer file sharing



Ciências
ULisboa

- Centralized, replicated index
- Decentralized edges




SETI@Home - Example




Ciências
ULisboa

*not quite a real
P2P system but
inspiring of
compute-
centric ones*




Peer-to-Peer Video Streaming



Topology of a Gnutella Network



Ciências
ULisboa



Source: Mihajlo A. Jovanovic, Fred S. Annexstein, and Kenneth A. Berman, Laboratory of Networks and Applied Graph Theory, University of Cincinnati.

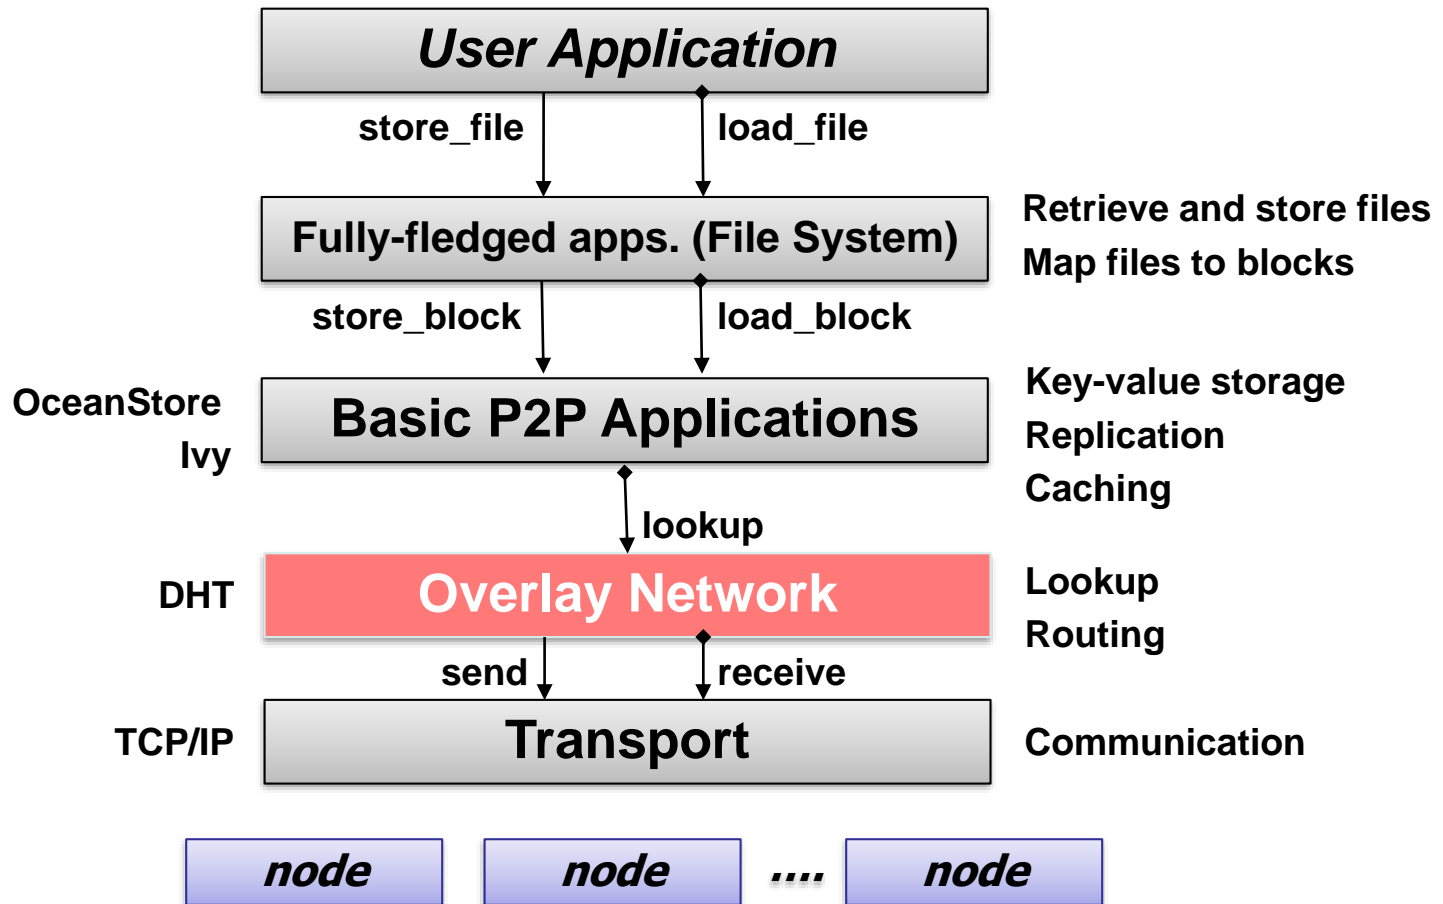
Peer-to-Peer System components

- Peer
 - An end-system, computer, end-user, application, etc.
- Overlay Network or Routing Overlay
 - Abstract (application-level) network maintained by nodes of a P2P system, keeping track of neighbours and/or a routing table
- P2P middleware
 - Intermediate layer for application-independent management of distributed P2P resources
- Leecher
 - A peer that uses and contributes resources (e.g. data)
- Seed
 - A peer that only contributes resources (e.g. data)

P2P system architecture



Ciências
ULisboa

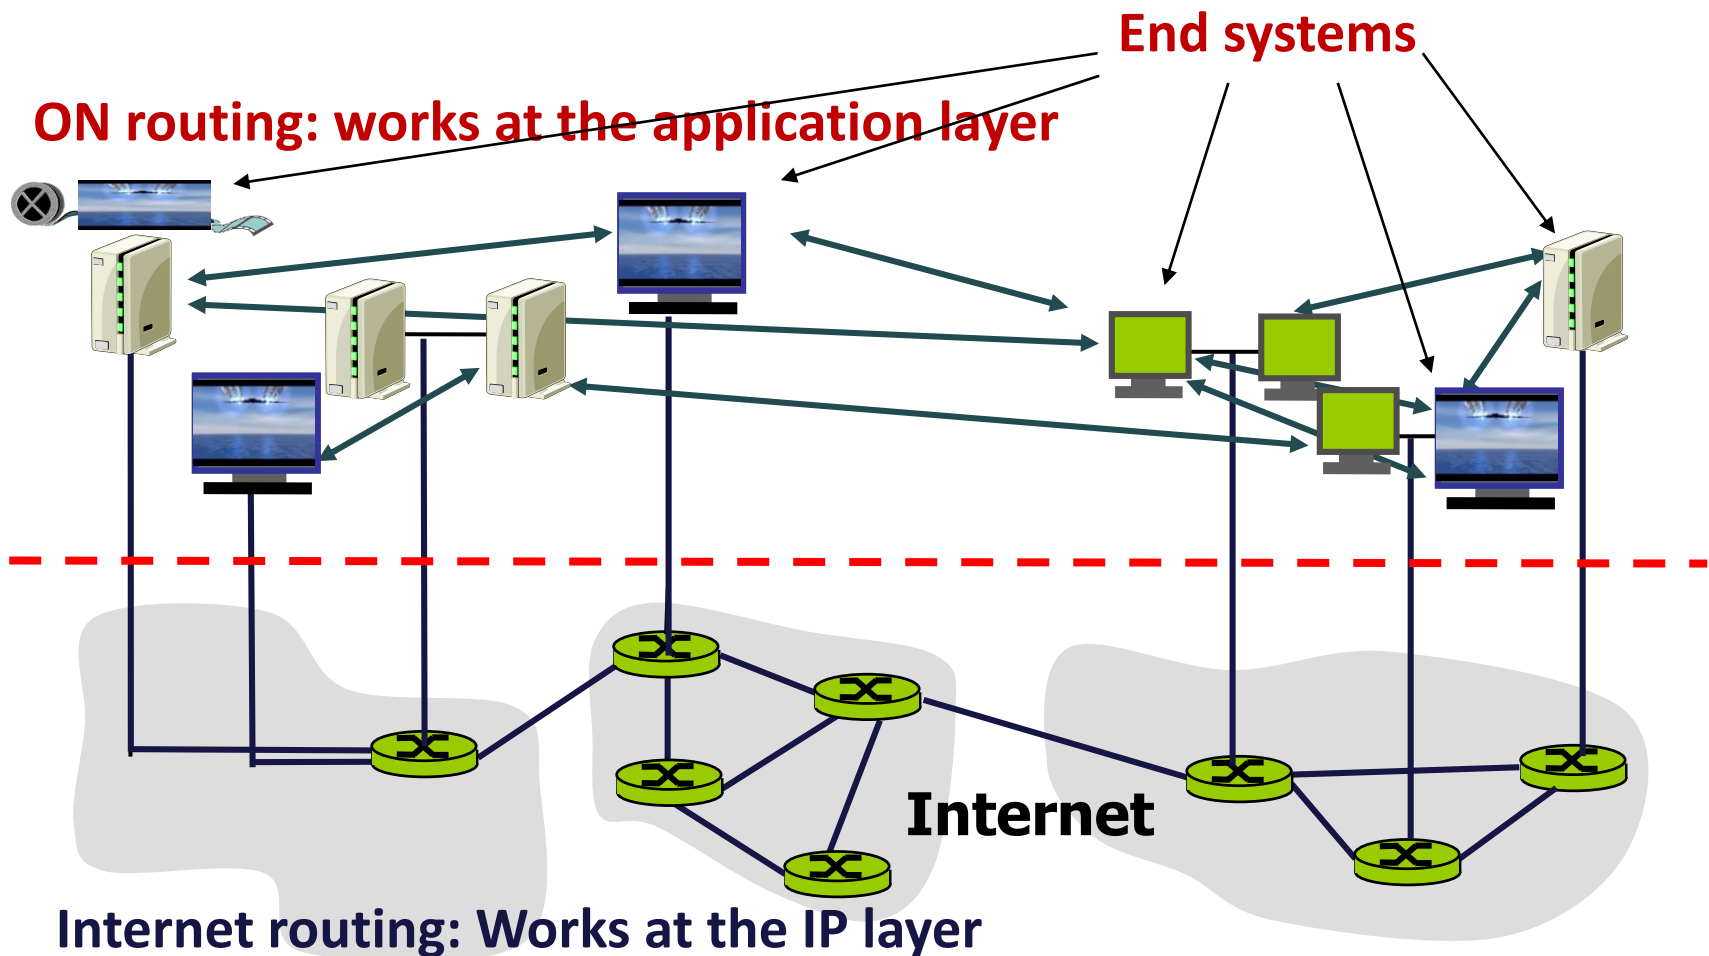


P2P systems vs. Overlay Networks

- A P2P system always resorts to an overlay network
- But an overlay network is not always a P2P system
- In P2P, an overlay network is formed by end-systems at application layer
- Overlay paths may be longer than IP paths (route “stretch”)

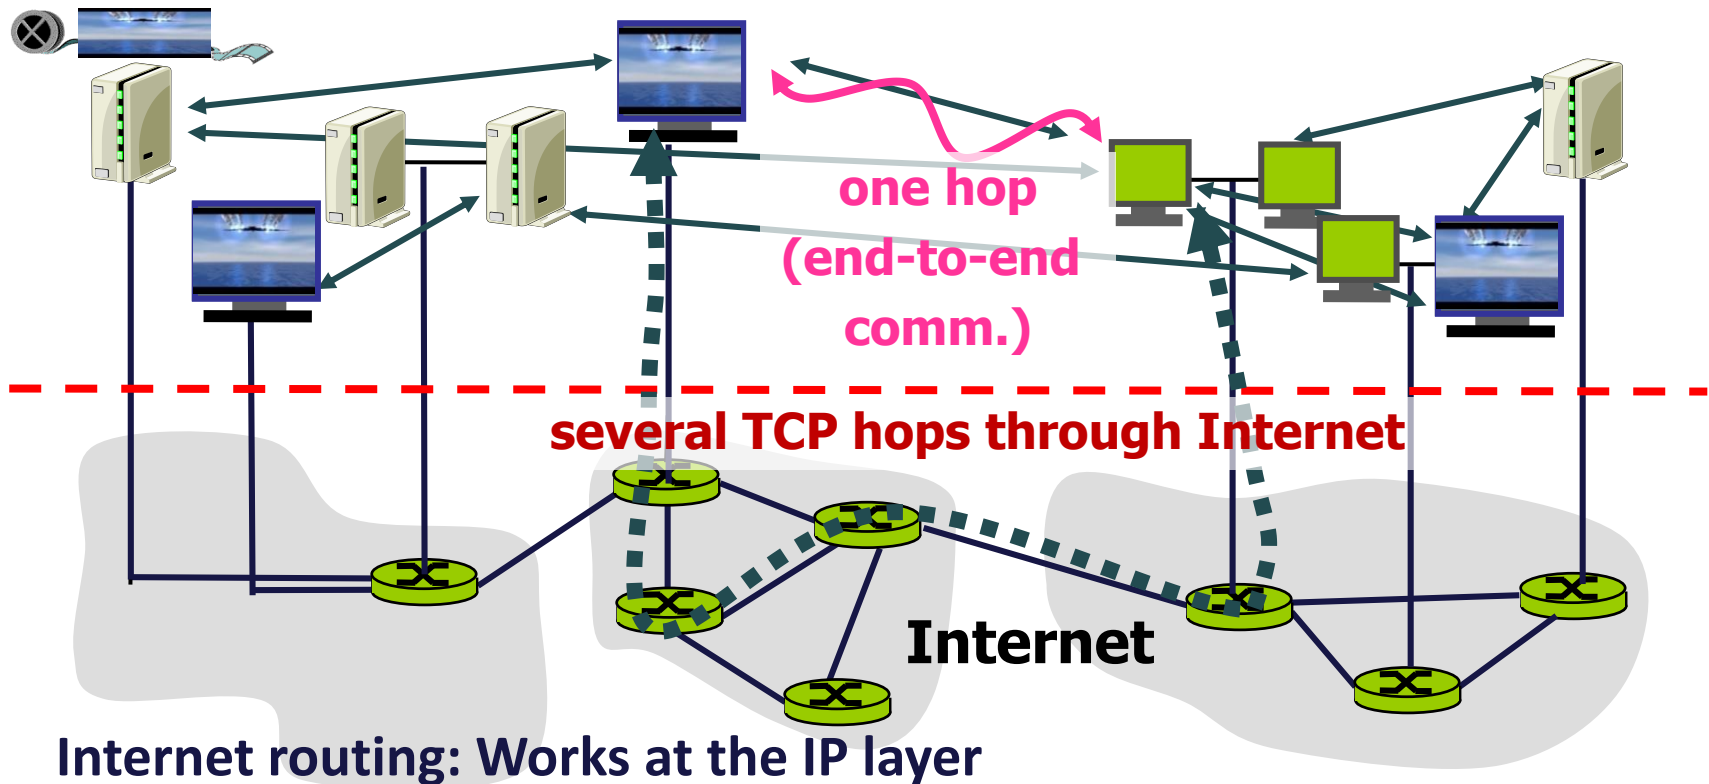
Routing stretch: the ratio between the ON distance (nr of ON hops) and the direct IP distance (nr of IP hops) of an ON route between x and y

P2P systems vs. Overlay Networks



P2P systems vs. Overlay Networks

ON routing: works at the application layer



Classification of P2P systems

- **Hybrid (also called centralized)** – Preserves some of the traditional C/S architecture. A central server links between clients, stores indices tables, etc
 - Napster
- **Unstructured** – no control over topology and resource (e.g. file) placement
 - Gnutella, Morpheus, Kazaa, etc
- **Structured** – topology is tightly controlled and placement of files is not random
 - Chord, CAN, Kademlia, PAST, Pastry, Tapestry, Tornado, BitTorrent, etc

Unstructured vs Structured P2P

Structured

PRO:

- Guaranteed to locate objects (assuming they exist);**
- Can offer time and complexity bounds on this operation;**
- Relatively low message overhead**

CON:

- Need to set up and maintain often complex overlay structures, which can be difficult; and costly to advertise/ discover resources;**
- Gets worse in highly dynamic environments (churn and load)**

Unstructured

- Self-organizing and naturally resilient to node failure**

- Only offers probabilistic guarantees on locating objects;**
- Excessive messaging overhead can affect scalability**