

---

# Race Conditions

Ibéria Medeiros  
Departamento de Informática  
Faculdade de Ciências da Universidade de Lisboa

# Linux Kernel "Dirty COW" Vulner. (Oct 2016)

---

"Linux users are being urged to patch servers to fix a vulnerability known as "**Dirty COW**." The **privilege elevation flaw** is caused by "**a race condition** ... in the way the Linux kernel's memory subsystem handles the copy-on-write (COW) breakage of private read-only memory mappings." Linux vendor Red Hat says that the flaw is being **actively exploited**. The flaw has been present in the **Linux kernel since 2007**, and is **trivial to exploit**. Linus Torvalds acknowledged that he **tried, unsuccessfully, to fix the problem more than 10 years ago.**"

# Race conditions (I)

---

- Violation of an assumption of atomicity
  - ☞ during a **window of vulnerability** or window of opportunity or window of inopportunity
  - ☞ 2 (or more) entities access concurrently the same object
- The vulnerability is always due to a problem of concurrency / lack of proper synchronization
  - ☞ between a target and malicious process(es), **or**
  - ☞ between several processes/threads of the target
- To exploit this kind of vulnerability, the attacker paces to break the assumption during the window of vulnerability

# Race conditions (II)

---

- Example: a service that generates unique sequential numbers
  - ☞ called by several threads concurrently
  - ☞ vulnerability allows returning the same number twice

```
int count = 0;    // global and shared by 2 threads
```

```
// Two threads execute this code concurrently
```

```
int getticket() {  
    count++;  
    return count;  
}
```



*There is an assumption of atomicity,  
which creates a window of vulnerability!*

To address this problem, the access to getticket() needs to be synchronized using locks, semaphores, ... or some other mechanism offered by the language.

# Race conditions (III)

---

- Sources of races
  - ➡ shared data: files and memory
  - ➡ preemptive routines (signal handlers)
  - ➡ multi-threaded programs
- We are going to study them in three scenarios
  - ➡ TOCTOU (*Time-Of-Check to Time-Of-Use*)
  - ➡ Temporary files
  - ➡ Internal concurrency and reentrant functions

---

**TOCTOU:**  
**TIME OF CHECK - TIME OF USE**

# TOCTOU Time-of-check to time-of-use

aka TOCTTOU

also symlink attack

- Typical case:

- ☞ A program running setuid root is asked to write to a file owned by the user running the program
- ☞ Root can write to any file so the program has to check if the actual user has the right to write to the file

*access() checks if the **Real UID** has write access to the file*

```
// 0 if the user has write privilege
if(!access(file, W_OK)) {
    f = fopen(file, "w+");
    write_to_file(f);
}
else {
    fprintf(stderr, "Permission denied\n");
}
```

} window  
vulnerab

Run this until success:

```
$ touch dummy
```

```
$ ln -s dummy link
```

```
$ program link &
```

```
$ rm link;\
```

```
ln -s /etc/passwd link
```

# access()

---

- `int access(const char *pathname, int mode);`
  - ☞ checks whether the calling process can access the file pathname according with mode
  - ☞ designed for setuid programs, does privilege check using the process' **real UID** instead of the **effective UID**
- However it is usually vulnerable to race conditions / TOCTOU
  - ☞ the file it checks can be altered before it is used
  - ☞ why? because it is identified by a path, so the object in the end can change
- Should never be used!



# A real example (I)

---


- Broken **passwd** command (SunOS, HP/UX)
  - ☞ it allows for instance to change the password
- **passwd** takes as parameter the password file to manipulate, then:
  1. Opens password file, retrieves user's entry, and closes file
  2. Create and open temporary file (*ptmp*) in the same directory as password file
  3. Open password file again, copy content to *ptmp* and update modified info
  4. Close both files, rename *ptmp* to be the password file

window  
vulnerab

# A real example (II)

- Attack script (interleaved with desired passwd execution):

```
$ mkdir evil
$ cp /etc/passwd evil/passwd
$ echo "hacker::0:0:::/bin/bash" >> evil/passwd
$ ln -s /etc link
$ passwd link/passwd
```



**passwd** step 1: open `link/passwd`, get user entry, close file

**passwd** step 2: create `ptmp`

```
$ rm link ; ln -s evil link
```

**passwd** step 3: open and copy `link/passwd` to `ptmp`, update modified info

```
$ rm link ; ln -s /etc link
```

**passwd** step 4: close files; move `ptmp` to `link/passwd` (i.e. `/etc/passwd`)

- `/etc/passwd` has a new user `hacker` with superuser privileges*

# Putting it to work

---

- Running *passwd* and exploiting the window of vulnerability seems hard...
- Solutions:
  - ☞ make a script that tries as many times as needed until the attack works
  - ☞ delay the program somehow
  - ☞ insert some random delays between operations
- Easier if the file system is distributed (e.g., NFS)

# stat() function family

---

- Collect information about files and symbolic links
  - 👉 `int stat(const char *pathname, struct stat *buf);`
    - pathname – file to be checked
    - `buf` – structure to be filled with data about the file
  - 👉 `lstat()` – the same but data returned is about the link itself, if `pathname` is a link
- These calls give info about the file
  - 👉 Owner, owning group, number of hardlinks to the file
  - 👉 Type of file (regular, dir, char device, block device, named pipe, symbolic link, socket)
- Apparently, these functions could be helpful to distinguish real files from links ...

# Vulnerability with lstat()

---

- setuid program has to open a file but does not want to be tricked into opening a symbolic link
  - ☞ trick: use lstat() instead of access()
- Vulnerable code in Kerberos 4 lib used in login daemon

```
if (lstat(file, &statb) != 0) goto out;  
if (!(statb.st_more & S_ISREG) || ...) goto out;  
if ((fd=open(file, O_RDWR|O_SYNC, 0)) < 0) goto out;
```

If it is **not** a regular file then exit.

- If file is not a link it is safe to open...  
but what if it changes after *lstat()* and before *open()* ?

# Vulnerability with lstat (cont)

---

- Is it still vulnerable if done in the opposite order?

```
fd = open(fname, O_RDONLY);  
if (fd == -1) perror("open");  
if (lstat(fname, &statb) != 0) die("file not  
there");  
if (!S_ISREG(statb.st_mode)) die("it's a symlink");
```

- Attacker creates link, then deletes it after the *open()* and before the *lstat()*

# Preventing file race conditions (I)

---

- Most file races have to do with pathnames
- When a call with a pathname is done (open, access, stat, lstat,...), the pathname is resolved until the inode is found
- If two calls are made one after the other, the path can lead to different *inodes*
- *So the protection is to avoid the two sequential resolutions and use filenames only when strictly necessary inside the program*

# Preventing file race conditions (II)

---

- Correct example – with file descriptors

```
fd=open("/tmp/bob", O_RDWR);  
fstat(fd, &sb);
```

- If someone unlinks and re-links /tmp/bob between the two calls, `fd` would still point to the same inode
- Unsafe: access, stat, lstat, chmod, chown
- Safe: fstat, fchmod, fchown



# Permission races

---

- Are we vulnerable in this case?

```
FILE *fp; //stream
int fd;
if (!(fp = fopen("myfile.txt", "w+"))) die("fopen");
fd = fileno(fp); // returns fd associated to stream
if (fchmod(fd, 0600)<0) // fchmod() to prevent race
    die("fchmod");
```

- If file does not exist

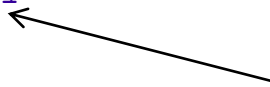
- ☞ *fopen()* creates it by calling *open()* with default permissions 0666
  - maybe lower as it is constrained by umask, but umask is inherited by the program
- ☞ Program changes it to 0600 but it is too late, a race is possible
- ☞ Attacker does not have to write in the file during the window, only open it for read/write (for instance)
- ☞ *Solution: set umask*

# Memory Races

---

- The following code is from a device driver in Windows 8.
- Are we vulnerable in this case?

```
PDWORD BufferSize = /* defined by a process in user mode */ ;
PUCHAR BufferPtr = /* defined by a process in user mode */ ;
PUCHAR LocalBuffer;
if (BufferSize > MAXSIZE) goto out;
LocalBuffer = ExAllocatePool(PagedPool, *BufferSize);
if (LocalBuffer != NULL) {
    RtlCopyMemory(LocalBuffer, BufferPtr, *BufferSize);
}
else goto out;
```



Copies an area of memory

- If the user changes the size of BufferSize between the check and its use, then he can write in areas of memory that were NOT supposed to (i.e., he can create a buffer overflow)

---

# TEMPORARY FILES

# Temporary files

---

- Have the same problems as others plus those derived from usually being in a shared dir

☞ */tmp; /var/tmp*

- Typical attack

1. Privileged program checks that there is no file X in */tmp*
2. Attacker races to create a link called X to some file, say */etc/passwd*
3. Privileged program attempts to create X and opens the attacker's file doing something undesirable that its privileges allow ...

- Two characteristics make the attack possible

☞ the **filename can be predicted** by the attacker

☞ the **race condition** between the **check** and the **file creation**

# mktemp(), tmpnam(),...

---

- Trying to address the problem: *mktemp()* creates unique, currently unused, filename from a template

```
strcpy(temp, "/tmp/tmpXXXXXX");  
if (!mktemp(temp)) die("mktemp");  
fd = open(temp, O_CREAT | O_RDWR, 0700);
```

- 👉 after *mktemp()* and before *open()* an attacker can link filename to some file
- 👉 *open()* would then not create but open an existing file

- tmpnam()* and *tempnam()* are similar

# Possible solutions

---

- Non-solutions

- ☞ Random file names: often the attacker can try the race many times, and if the name generator can be predicted ...
- ☞ Locks: many implementations are enforced by convention, not mandatory

- Acceptable solutions

- ☞ Use long random number (e.g. 64 bits), set umask appropriately (e.g. 0066), and open the file
- ☞ Use the safe calls → see *next slide*

# mkstemp(), tmpfile(),...

---

- Possible solutions:

*int mkstemp(char \*template)*

much safer than *mktemp()* since it checks for uniqueness, creates and opens with rw privileges only for the owner (0600)

*FILE \*tmpfile(void)*

is similar, but in the past sometimes it was implemented on top of *mktemp()*, and therefore could be vulnerable to attacks

*char \*mkdtemp(char \*template)*

is similar but creates a directory with permissions 0700

---

# CONCURRENCY AND REENTRANT FUNCTIONS



# Concurrency

---

- In the previous cases, concurrency **is (mainly) created by the attacker**, with malicious intention
- In programs there is also the normal concurrency of operations that need to **access shared objects**
- Operations may have to be executed atomically, i.e., without interruption
- Mutual exclusion is a solution for some of the problems
  - ☞ mutexs, semaphores, transactions, etc.
- Difficulties with mutual exclusion
  - ☞ **starvation**: a thread is never scheduled for execution
  - ☞ **deadlock**: threads inter-block themselves

# Reentrancy (I)

---

- A function is reentrant if several instances of the function can be executed in parallel in the same address space
  - ☞ for example, if it works correctly even if a thread is interrupted by another thread that calls the same function
- ☞ Example:
  - non-reentrant function supposed to give unique ticket

```
int count=0;    //global var
int getticket() {
    count++;
    return count;
}
```

} atomic?

# Reentrancy (II)

---

- Separating the reentrancy in two main sources of problems
  - ☞ Thread-safety – reentrancy in relation to several threads
  - ☞ Async-signal-safety – reentrancy in relation to signals
- Conditions for a function to be reentrant
  - ☞ does not use: static variables, global variables, other shared resources like libraries (i.e., uses only local non-static variables and function parameters)
  - ☞ only calls reentrant functions (namely libraries)
- Notice: conditions are **sufficient but not necessary!**
  - ☞ the use of global vars may **not** prevent the function from being reentrant

# Signal handlers

---

- Signals are a Unix/POSIX mechanism to indicate asynchronous events to a process
  - ☞ often the semantics varies in different Unix flavors
- Can be treated by a signal handler (a function)
  - ☞ or ignored or blocked (except SIGKILL and SIGSTOP)
- Signal handlers have to be **asynchronous-safe** (or **async-safe** or **signal-safe**):
  - ☞ have to run correctly even if interrupted by other asynchronous events (i.e., by a signal) **and**
  - ☞ they should not corrupt the part of the program that was interrupted
- Otherwise they may be vulnerable (or create vulnerabilities) and often are!

# Example of Signal Vulnerability

```
int size;
FILE *fd;
Char *bglobal;
int main(int argc, char **argv){
    signal(SIGTERM, complete);
    signal(SIGINT, complete);
    signal(SIGSEGV, complete);
    // open file for input
    fd = fopen("input.txt", "r");
    // now start two threads →
}
void thread_readData(){
    bglobal = malloc(MAX_BYTES);
    while (true) {
        wait_workDone(); // from thrd2
        size = fread(bglobal, 1,
                     MAX_BYTES, fd);
        more_work(); // warn thrd2
        printf("Read %d bytes.\n", size);
    }
}
```

```
void thread_workData(){
    char buffer[MAX_BYTES];
    while (true){
        wait_work(); //wait thread1
        memcpy(buffer, bglobal, size);
        done_work(); // warn thread1
        // does some processing
    }
}
int complete(int sig){
    printf("recv signal. Done!.\n");
    free(bglobal);
    fclose(fd);
    exit(1);
}
```

If the signal is executed :  
What if **complete** is called two times?  
What if **signal occurs** before memory is allocated or the file is open?  
What if **signal occurs** when printf() was being executed?

# Signals

---

- Some solutions

- ☞ Read very carefully the man pages!!!
- ☞ Make signal handler simple; ideally only set a flag
- ☞ Use system calls safe for signal handlers
- ☞ Block signals
  - (1) inside signal handlers and
  - (2) during non-atomic operations in the program

- Keep in mind that

- ☞ mutual exclusion mechanisms typically cannot be used because they can lead to deadlocks (depending on how signals are configured)
- ☞ alternatively, it is possible to use mechanisms that temporarily block the delivery of signals

# Servlets

---

- Usually
  - ☞ Called by several threads concurrently
  - ☞ Each servlet is only one object in memory...  
... so class attributes are shared resources (like global vars in C)
- Servlet that can give two users the same count

```
public class Counter extends HttpServlet {  
    int count = 0;    //shared!!  
    public void doGet(..., HttpServletResponse out)  
        throws... {  
        PrintWriter p = out.getWriter();  
        count++;  
        p.println(count + " hits so far!");  
    }  
}
```

*There is an assumption of atomicity,  
which creates a window of vulnerability!*

# Solutions

---

```
public class Counter extends HttpServlet {
    int count = 0;    //shared!!
    public synchronized void doGet(..., HttpServletResponse out)
        throws... {
        PrintWriter p = out.getWriter();
        count++;
        p.println(count + " hits so far!");
    }
}
```

not efficient (if method longer...)

---

```
public class Counter extends HttpServlet {
    int count = 0;    //shared!!
    public void doGet(..., HttpServletResponse out) throws... {
        int my_count;    // NOT shared!!!
        PrintWriter p = out.getWriter();
        synchronized(this) {    my_count = ++count;    }
        p.println(my_count + " hits so far!");
    }
}
```



# Bibliography

---

- M. Correia, P. Sousa, Segurança no Software, FCA Editora, 2017 (see chapter 6)

Other references: