

5. Failure Detectors and Consensus

Local Failure Detection

- i.e., detection of failed processes from inside a host
- Detection can be considered reliable
 - Why?
- Various mechanisms:
 - Self-checking routines (e.g. parity checks)
 - Guardian components (e.g. memory access checkers)
 - Watch-dogs (hardware-based, software-based)
- Temporal problems may still affect accuracy
 - Why?
- *(from now on: distributed failure detection)*
- *(in this lecture: failure = crash)*

Basic Failure Detectors: Ping

```
Alyssons-MBP:~ anbessani$ ping www.google.com
PING www.google.com (172.217.16.228): 56 data bytes
64 bytes from 172.217.16.228: icmp_seq=0 ttl=56 time=12.662 ms
64 bytes from 172.217.16.228: icmp_seq=1 ttl=56 time=13.100 ms
64 bytes from 172.217.16.228: icmp_seq=2 ttl=56 time=13.238 ms
64 bytes from 172.217.16.228: icmp_seq=3 ttl=56 time=13.011 ms
64 bytes from 172.217.16.228: icmp_seq=4 ttl=56 time=12.356 ms
64 bytes from 172.217.16.228: icmp_seq=5 ttl=56 time=13.119 ms
64 bytes from 172.217.16.228: icmp_seq=6 ttl=56 time=14.157 ms
64 bytes from 172.217.16.228: icmp_seq=7 ttl=56 time=13.245 ms
64 bytes from 172.217.16.228: icmp_seq=8 ttl=56 time=72.393 ms
^C
--- www.google.com ping statistics ---
9 packets transmitted, 9 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 12.356/19.698/72.393/18.636 ms
Alyssons-MBP:~ anbessani$
```

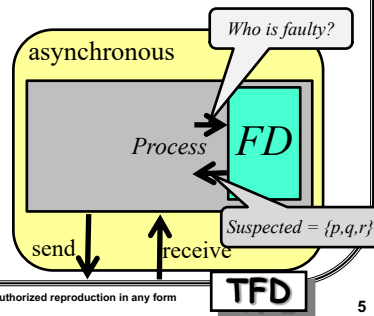
- Typically ICMP: echo-request, echo-reply

Engineering Failure Detection

- Heartbeat mechanism:
 - Observed component periodically sends messages
 - Usually called “I’m alive” messages or “heartbeats”
 - When these messages are not received, the component is considered faulty
- Probe mechanism (similar to a ping):
 - Observed component waits for a probe message and replies
 - Probe comes from the failure detector
 - Component replies with “I’m alive” message
- Which one is better?

Failure Detectors (FDs)

- Ping provides suspicions (**hints**) of failures (crashes)
- Conceptually, FDs are **distributed oracles** that provide hints about process failures
 - local modules attached to processes
 - modules may communicate using the network
- Using FDs allow the specification of algorithms that does not express timing assumptions
 - System can be asynchronous
 - (no time assumptions)
 - But failures can be detected
 - (through FD module)
 - Failure detection is **modularized**
 - Separation of concerns!



© 2002-2018 A. Bessani & A. Casimiro & M. Correia, All rights reserved, no unauthorized reproduction in any form

5

Recalling Safety and Liveness

- **Safety**
 - The measure in which a service/program does not do bad things
 - Safety properties specify that wrong events never take place
 - A safety property specifies that a predicate P is always true
- **Liveness**
 - The measure in which a service/program does good things
 - Liveness properties specify that good events eventually take place
 - A liveness property specifies that predicate P will eventually be true

(Recall the properties we used to define reliable broadcast, etc.)

© 2002-2018 A. Bessani & A. Casimiro & M. Correia, All rights reserved, no unauthorized reproduction in any form

TFD

6

Failure detection

properties and problems

- Consistency of distributed failure detection:
 - A fundamental problem in asynchronous distributed systems is **how to differentiate a slow process from a crash**
 - Ideally, when a process goes down all the other processes must know about it to coordinate their actions to take corrective actions
 - Such FD consistency can be required for **safety** or **liveness**
 - What's the difference?
- One of the biggest issues when designing a practical distributed system is if we can reliably detect process failures
 - And, if not, how we deal with misdetections?

Failure detection

properties and problems

Properties of Consistent Failure Detection

- Strong Accuracy
 - A **safety** requirement, specifying that no correct process is ever considered faulty
- Strong Completeness
 - A **liveness** requirement, specifying that a failure must be eventually detected by every correct process

Is it possible to satisfy these properties on the internet?

Failure detection

properties and problems

- On the possibility of implementing strong accuracy and completeness...
- Assuming that message processing requires negligible time, processes have clocks with bounded drift, and “perfect channels” are available, heartbeat/probe exchanges meet strong accuracy and completeness
 - Perfect channel = synchronous + reliable
- Such a detector is called perfect failure detector:
 - If a process crashes, all correct processes will note the absence of the heartbeat and detect the failure
 - No correct process is ever suspected

Failure detection

properties and problems

- What happens when the channels are not perfect?
 - Either imperfection can be fixed by some simple protocol
 - Or imperfection is impossible to overcome
- Communication channel is not perfect but has some mild imperfection:
 - For instance, it makes at most k consecutive omissions
 - Solution: transform the imperfect channel into a perfect channel
 - E.g., each heartbeat can be retransmitted $k + 1$ times, effectively ensuring that it is observed by all correct processes.

Failure detection

properties and problems

- Channel imperfection impossible to overcome:
 - Lack of bounds on the number and type of faults that may happen (e.g., number of channel omission faults not bounded)
- Consequence:
 - Now, if a process does not receive any heartbeat message from the other process this may have two causes:
 - Because the other process is failed or
 - Because the channel has dropped all heartbeats sent so far!!!

Failure detection

properties and problems

- Channel imperfection impossible to overcome:
 - Lack of bounds for the timely behavior of system components (processes or links) – *asynchrony*
- Consequence:
 - No way to distinguish *missing* from “*extremely slow*” heartbeats
 - Happens if a link can delay a message arbitrarily, or if a process can take an arbitrary amount of time to make a processing step
 - Perfect failure detection cannot be implemented in asynchronous systems!!!

Hierarchy of FDs

- Chandra and Toueg (1992 and 1996) have defined the notion of **unreliable FD** and a hierarchy of such FDs
- Perfect failure detector \mathcal{P}
 - **Strong Accuracy**. No process is suspected before it crashes.
 - **Strong Completeness**. Eventually every process that crashes is permanently suspected by every correct process.

Failure detection properties and problems

These properties can be relaxed:

- **Weak Accuracy**
 - At least one correct process is never suspected by all correct processes
- **Weak Completeness**
 - Eventually every process that crashes is permanently suspected by at least one correct process

Even Weak Accuracy is impossible in non-synchronous systems!

Therefore:

- **Eventual Strong Accuracy**
 - There is a time after which all correct processes are never considered failed by any correct processes
- **Eventual Weak Accuracy**
 - There is a time after which some correct process is never considered failed by any correct processes

Eventually weak FD: $\diamond W$

- The weakest FD of the Chandra-Toueg hierarchy is defined by the following properties:
- Eventual Weak Accuracy**
 - There is a time after which some correct process is never considered failed by any correct processes
- Weak Completeness**
 - Eventually every process that crashes is permanently suspected by at least one correct process
- Consensus is solvable in an asynchronous system with an eventually weak FD
 - Notice that this system model is stronger than the asynchronous system model; the eventually weak FD is not implementable in the latter!
 - It is impossible to solve consensus with a weaker detector, as proved by Chandra, Hadzilacos and Toueg (1996)

© 2002-2018 A. Bessani & A. Casimiro & M. Correia, All rights reserved, no unauthorized reproduction in any form

TFD

15

Hierarchy of FDs

- Note:**
 - the diamond stands for “eventually”
 - S is an important FD for “practical” synchronous systems

Completeness	Accuracy			
	Strong	Weak	Eventually strong	Eventually weak
Strong	Perfect \mathcal{P}	Strong \mathcal{S}	Eventually perfect $\diamond \mathcal{P}$	Eventually Strong $\diamond \mathcal{S}$
Weak	\mathcal{Q}	Weak \mathcal{W}	$\diamond \mathcal{Q}$	Eventually weak $\diamond \mathcal{W}$

© 2002-2018 A. Bessani & A. Casimiro & M. Correia, All rights reserved, no unauthorized reproduction in any form

TFD

16

Weakest FD revisited

Completeness	Accuracy			
	Strong	Weak	Eventually strong	Eventually weak
Strong	Perfect \mathcal{P}	Strong \mathcal{S}	Eventually perfect $\diamond \mathcal{P}$	Eventually Strong $\diamond \mathcal{S}$
Weak	\mathcal{Q}	Weak \mathcal{W}	$\diamond \mathcal{Q}$	Eventually weak $\diamond \mathcal{W}$

w/ crash faults it is trivial to obtain **Strong Completeness** from **Weak Completeness** (so using $\diamond \mathcal{S}$ is more or less the same as $\diamond \mathcal{W}$):

```

outputp ← ∅
cobegin
  || Task 1: repeat forever
    {p queries its local failure detector module  $\mathcal{D}_p$ }
    suspectsp ←  $\mathcal{D}_p$ 
    send (p, suspectsp) to all

  || Task 2: when receive (q, suspectsq) for some q
    outputp ← (outputp ∪ suspectsq) - {q}
coend

```

Here is a transformation

{output_p emulates \mathcal{D}_p }

Summary: FD classes and properties

Completeness	Accuracy			
	Strong	Weak	Eventually strong	Eventually weak
Strong	Perfect \mathcal{P}	Strong \mathcal{S}	Eventually perfect $\diamond \mathcal{P}$	Eventually Strong $\diamond \mathcal{S}$
Weak	\mathcal{Q}	Weak \mathcal{W}	$\diamond \mathcal{Q}$	Eventually weak $\diamond \mathcal{W}$

- **Strong Completeness.** Eventually every process that crashes is permanently suspected by every correct process.
- **Strong Accuracy.** No process is suspected before it crashes.
- **Weak Accuracy.** At least one correct process is never suspected by all correct processes
- **Eventual Strong Accuracy.** There is a time after which all correct processes are never considered failed by any correct processes
- **Eventual Weak Accuracy.** There is a time after which some correct process is never considered failed by any correct processes

Implementing a FD of class $\diamond P$

(that also implements weaker FDs: $\diamond S$ and $\diamond W$)

- Consider the following system model:
 - Partial synchrony
 - For every execution there is a Global Stabilization Time (GST) after which there are **bounds** on the relative process speeds and communication delays (both GST and the bounds are unknown)
 - Reliable channels
 - Processes can fail by crashing

Implementing a FD of class $\diamond P$

Every process p executes the following:

```

outputp ← ∅
for all  $q \in \Pi$ 
   $\Delta_p(q) \leftarrow$  default time-out interval
  { $\Delta_p(q)$  denotes the duration of  $p$ 's time-out interval for  $q$ }

cobegin
|| Task 1: repeat periodically
  send "p-is-alive" to all
|| Task 2: repeat periodically
  for all  $q \in \Pi$ 
    if  $q \notin \text{output}_p$  and
       $p$  did not receive "q-is-alive" during the last  $\Delta_p(q)$  ticks of  $p$ 's clock
      outputp ← outputp ∪ { $q$ }
      { $p$  times-out on  $q$ : it now suspects  $q$  has crashed}
|| Task 3: when receive "q-is-alive" for some  $q$ 
  if  $q \in \text{output}_p$ 
    outputp ← outputp - { $q$ }
     $\Delta_p(q) \leftarrow \Delta_p(q) + 1$ 
    { $p$  knows that it prematurely timed-out on  $q$ }
    {1.  $p$  repents on  $q$ , and}
    {2.  $p$  increases its time-out period for  $q$ }
coend
  
```

Solving consensus with S and $\diamond S$

- Consensus is impossible in asynchronous systems, but solvable with failure detectors
- Consensus:
 - Termination: Every correct process eventually decides on some value.
 - Validity: If a process decides v , then v was proposed by some process.
 - Agreement: No two correct processes decide differently.
- Let's study the protocol by Mostefaoui and Raynal (1999)

Solving consensus with S and $\diamond S$

Step by Step

- Step 0
 - Initially, all processes P_i have their own $est_i = v_i$ of decision
 - Some process must send a proposal
 - Let's use the **rotating coordinator** approach
 - All processes know n and the coordinator c for each round r
($c = r \bmod n$)
 - Coordinator reliably sends its est to all processes
 - All processes wait for receiving est from coordinator
 - When they receive est , they decide est
- Doesn't work. Why?
 - If the coordinator fails (by crashing), then **termination** is violated!!
 - Failure of the coordinator must be detected (and recovered)

Solving consensus with S and $\diamond S$

Step by Step

• Step 1: Proposal

- Let's use a Failure Detector to allow progress when c crashes
 - Consider a **Strong FD** or an **Eventually Strong FD**
 - If c crashes, all processes will eventually detect the failure (SC)
 - But it is possible that some process will not wait enough time
 - All processes record the est_from_c , received from c
 - At some point (the end of initial phase) est_from_c can be equal to v (c 's proposal), or can be empty (\perp)
- ```

(3) $c \leftarrow (r_i \bmod n) + 1$; $est_from_c_i \leftarrow \perp$; $r_i \leftarrow r_i + 1$; % round $r = r_i$ %
(4) case ($i = c$) then $est_from_c_i \leftarrow est_i$ and send it to all
(5) ($i \neq c$) then wait ((EST(r_i, v) is received from p_c) \wedge ($c \in suspected_i$));
(6) if (EST(r_i, v) has been received) then $est_from_c_i \leftarrow v$
(7) endcase; % $est_from_c_i = est_c$ or \perp %

```

## • It's not yet possible to decide! Why?

- How to ensure that all others will decide the same value?

Solving consensus with  $S$  and  $\diamond S$ 

## Step by Step

## • Step 2: Confirmation

- Let's make all processes send their  $est\_from\_c$  to each other
    - Transmission is reliable, but processes may fail
    - A process can wait forever for all (n) possible  $est\_from\_c$
    - A process finishes this phase when it has received  $est\_from\_c$  values from "enough" processes – a quorum (set Q)
    - "enough" = any two quorums must have a non-empty intersection
    - **The size of Q depends on the failure detector employed**
  - Received  $est\_from\_c$  values are stored in a  $rec$  (quorum) set
- ```

(8)   $\forall j$  do send EST( $r_i, est\_from\_c_i$ ) to  $p_j$  enddo;
(9)  wait until ( $\forall p_j \in Q_i$ : EST( $r_i, est\_from\_c$ ) has been received from  $p_j$ );
      %  $Q_i$  has to be a live and safe quorum %
      % For  $S$ :  $Q_i$  is such that  $Q_i \cup suspected_i = \Pi$  %
      % For  $\diamond S$ :  $Q_i$  is such that  $|Q_i| = \lceil (n+1)/2 \rceil$  %
(10) let  $rec_i = \{est\_from\_c \mid \text{EST}(r_i, est\_from\_c) \text{ is received at line 5 or 9}\};$ 
      %  $est\_from\_c = \perp$  or  $v$  with  $v = est_c$  %
      %  $rec_i = \{\perp\}$  or  $\{v\}$  or  $\{v, \perp\}$  %

```

Solving consensus with S and $\diamond S$

Step by Step

- How to achieve overlapping quorums?
 - Let's consider an FD of class S
 - For sure, one correct process is never suspected
 - The condition to finish phase 2 is therefore waiting for messages from all correct processes
 - Even if the FD suspects all processes but one, this one will be included in all quorums
 - This means that the algorithm works for $f < n$
 - What if the FD of $\diamond S$ class?
 - FD can commit an arbitrary number of mistakes in a round
 - It is necessary to wait for messages of a majority of processes to ensure that any two **quorums** have at least one common process
 - The output of the FD (list of suspects) is not used in this case
 - This algorithm works because it is assumed that $f < n/2$

Solving consensus with S and $\diamond S$

Step by Step

- Step 2: Confirmation (cont.)
 - The **rec** set of any process can contain only **est_from_c** = v or \perp
 - Given these values, is it now possible to decide?
 - Let's see the three possible outcomes for the **rec** set:
 1. $rec_i = \{v\} \Rightarrow (\forall p_j : (rec_j = \{v\}) \vee (rec_j = \{v, \perp\}))$
 2. $rec_i = \{\perp\} \Rightarrow (\forall p_j : (rec_j = \{\perp\}) \vee (rec_j = \{v, \perp\}))$
 3. $rec_i = \{v, \perp\} \Rightarrow (\forall p_j : (rec_j = \{v\}) \vee (rec_j = \{\perp\}) \vee (rec_j = \{v, \perp\}))$
 - 1. P_i can decide v , since any other process also decides v , or will take v as its own **est** for the next round
 - 2. P_i will not decide, and no other process will decide
 - 3. P_i will not decide, but since other processes might have decided v , it will update its own **est** to v

Solving consensus with S and $\diamond S$

Step by Step

• Step 3: Decision

- Just before deciding (returning **est**) the decision is sent to all processes (reliably broadcast)
 - Since the deciding process will stop executing the algorithm, this ensures that all correct processes decide, in case they have not decided in the same round
 - If the process crashes before successfully completing R-Broadcast, this may only delay decision (on same value) of correct processes

```

(11) case ( $rec_i = \{\perp\}$ ) then skip
(12)      ( $rec_i = \{v\}$ ) then  $\forall j \neq i$  do send DECIDE( $v$ ) to  $p_j$  enddo; return( $v$ )
(13)      ( $rec_i = \{v, \perp\}$ ) then  $est_i \leftarrow v$ 
(14) endcase
(15) enddo

(16) task T2: upon reception of DECIDE( $v$ ):
       $\forall j \neq i$  do send DECIDE( $v$ ) to  $p_j$  enddo; return( $v$ )

```

Function Consensus(v_i)

cobegin

```

(1) task T1:  $r_i \leftarrow 0$ ;  $est_i \leftarrow v_i$ ; %  $v_i \neq \perp$  %
(2) while true do
(3)    $c \leftarrow (r_i \bmod n) + 1$ ;  $est\_from\_c \leftarrow \perp$ ;  $r_i \leftarrow r_i + 1$ ; % round  $r = r_i$  %
(4)   case ( $i = c$ ) then  $est\_from\_c \leftarrow est_i$ 
(5)     ( $i \neq c$ ) then wait ((EST( $r_i, v$ ) is received from  $p_c$ )  $\vee$  ( $c \in suspected_i$ ));
(6)       if (EST( $r_i, v$ ) has been received) then  $est\_from\_c \leftarrow v$ 
(7)   endcase; %  $est\_from\_c = est_c$  or  $\perp$  %
(8)    $\forall j$  do send EST( $r_i, est\_from\_c$ ) to  $p_j$  enddo;
(9)   wait until ( $\forall p_j \in Q_i$ : EST( $r_i, est\_from\_c$ ) has been received from  $p_j$ );
      %  $Q_i$  has to be a live and safe quorum %
      % For  $S$ :  $Q_i$  is such that  $Q_i \cup suspected_i = \Pi$  %
      % For  $\diamond S$ :  $Q_i$  is such that  $|Q_i| = \lceil (n+1)/2 \rceil$  %
(10)  let  $rec_i = \{est\_from\_c \mid \text{EST}(r_i, est\_from\_c) \text{ is received at line 5 or 9}\}$ ;
      %  $est\_from\_c = \perp$  or  $v$  with  $v = est_c$  %
      %  $rec_i = \{\perp\}$  or  $\{v\}$  or  $\{v, \perp\}$  %
(11)  case ( $rec_i = \{\perp\}$ ) then skip
(12)    ( $rec_i = \{v\}$ ) then  $\forall j \neq i$  do send DECIDE( $v$ ) to  $p_j$  enddo; return( $v$ )
(13)    ( $rec_i = \{v, \perp\}$ ) then  $est_i \leftarrow v$ 
(14)  endcase
(15) enddo

```


```

(16) task T2: upon reception of DECIDE( $v$ ):
       $\forall j \neq i$  do send DECIDE( $v$ ) to  $p_j$  enddo; return( $v$ )

```

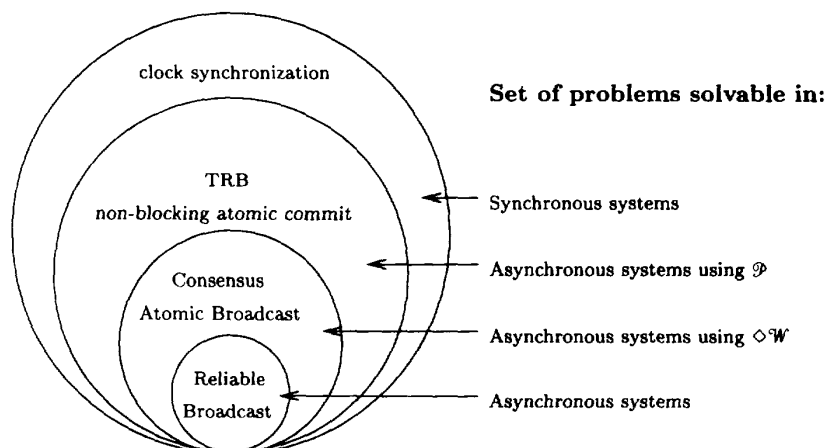
coend

Solving consensus with S and $\diamond S$



- A note about minimal number of communication steps:
 - With minimal number of processes ($n > 2f$ or $n > f$), consensus can be solved in two communication steps
 - Notice that RAFT, VSR and Paxos require three steps
 - With more processes ($n > 3f$), consensus can be solved in one communication step

What problems are solvable with (classical) failure detectors?



Bibliography

- T. Chandra and S. Toueg, **Unreliable Failure Detectors for Reliable Distributed Systems**. Journal of the ACM, 43(2):225-267, 1996.
- T. Chandra, V. Hadzilacos and S. Toueg, **The Weakest Failure Detector for Solving Consensus**. Journal of the ACM, 43(4):685–722, 1996.
- A. Mostéfaoui and M. Raynal. **Solving Consensus Using Chandra-Toueg's Unreliable Failure Detectors: A General Quorum-Based Approach**, DISC, 1999.