


8. Recovery

Replica Recovery

- Consider a replicated service in which replicas can crash and recover (crash-recovery fault model)
- Recovery from crashes requires that the recovering replica obtains current state from other replicas
- Without stable storage:
 - Cooperative recovery from other replica(s) without stable storage
 - **Complete state transfer** can make recovery very slow
 - Lose state if all replicas fail
- With stable storage:
 - Recover some past state before failure
 - Cooperative recovery of missing commands from other replica(s)
 - Execute commands from the log until current state
 - State recovery is much faster

Recovery with state transfer



- Recovering replica (p3) starts by resuming communication with the replica set, e.g. if the set was using some form of group communication
- It starts receiving all messages, but still discards them
- Next, sends a request to join the replica group, delivered in total order to all replicas, including the joining replica, marking a **cut S_j** in the global system state
- The join request triggers a state-transfer operation (after join is completed)
- p2 checkpoint its state at this point (S_j), and sends it to p3
- p3 starts logging any messages that arrive after the cut S_j
- New requests (m_k) can continue to be processed by all replicas except p3

Checkpoint-based Rollback-recovery

(Scope and basics)

- Rollback-recovery is mostly suited for **long running applications**
 - Scientific applications (simulation, optimization)
 - Telecommunication applications
- Used in **local and distributed computations**
- Requires **stable storage** that survives failures
- State information is periodically saved (**checkpoints**)
- Other information may also be saved (**logs**)
- After a failure, the saved information can be used to restart the computation from an intermediate state

Issues in checkpointing

- At what level should we checkpoint?
 - Should it be made transparent to the user?
 - What are the trade-offs involved in the decision?
- How many checkpoints should be done?
 - Frequently?
 - Sparingly?
- When should they be done?
 - Upon every event, or only when some events occur?
- How to minimize checkpoint overhead?
- How to do distributed checkpoints?
 - Coordinated?
 - Uncoordinated?

Checkpointing approaches

- **Uncoordinated checkpoints**
 - Every process takes a checkpoint independently from other processes
 - Since there is no coordination, it may be possible that checkpoints are not consistent across system (**domino effect**)
- **Coordinated checkpoints**
 - Processes coordinate to meet a **consistent global state** before taking the checkpoint
 - Only the last checkpoint needs to be kept
- **Communication induced checkpoints**
 - Processes take checkpoints independently
 - But they select when to take a checkpoint, to avoid the domino effect (see next slide)


Domino effect

(Message-passing systems)

- What is the domino effect?
 - Rollback of one process meets an inconsistent global state, forces rollback of another process, which in turn forces first process to rollback again, and so on
- This can only happen with independent checkpointing
- The solution is to use coordinated checkpointing
 - A system-wide consistent state is saved
 - The Chandy-Lamport algorithm is a widely used solution

Domino effect

(illustrated)



1. p_1 fails, and then recovers, rolling back to checkpoint C_a
2. Evidence of sending message m_i no longer exists
3. So, p_2 is forced to rollback to checkpoint C_b
4. However, this "unsends" message m_j and p_3 is forced back to C_c
5. Rollback propagation will bring system back to initial state

Logging

- Checkpointing allows recovering from a past state
- If done frequently:
 - No need to go too far back in state
 - Redoing undone actions will be faster
- but
 - Overhead of checkpointing can become unbearable
 - The problem is amplified if state is large
- Solution: **reduce frequency of checkpoints!**
- Unfortunately this may not be possible if some actions
 - Leave traces outside the system
 - Are not deterministic

Log-based rollback-recovery

- The solution to 1) avoid frequent checkpoints and 2) ensure deterministic recovery is to **log all non-deterministic events**
- The state can then be reconstructed from the most recent checkpoint + log
 - All the non-deterministic events are replayed in their exact order
- Major approaches to log-based rollback-recovery:
 - **Pessimistic logging**: every event is logged before taking effect
 - **Optimistic logging**: asynchronous logs, which may be lost
 - **Causal logging**: keeps track of causal relations among events

Log-based rollback-recovery

- Assumes that non-deterministic events can be logged
 - E.g., Received messages or internal events
 - Sent messages are deterministic events!
 - The execution is a sequence of deterministic state intervals
- Some checkpoints are taken to avoid long roll-backs
- After failure the state is recovered using checkpoints and logged information of non-deterministic events
- After recovery there will be no orphan processes
 - No process will be in a state that depends on a non-reproducible non-deterministic event

Log-based rollback-recovery (Protocols)

- Pessimistic log-based rollback-recovery
 - Guarantee that orphans are not created
 - Simple recovery, garbage collection and output commit
 - High performance overhead
- Optimistic log-based rollback-recovery
 - Reduced performance overhead
 - Orphans may be created
 - Recovery, garbage collection and output commit may be harder
- Causal log-based rollback-recovery
 - Guarantee that orphans are not created
 - Reduced performance overhead and fast output commit
 - Complex recovery and garbage collection