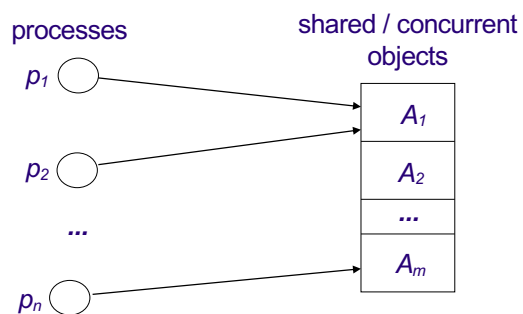


11. Shared Memory Objects

Introduction

- **Concurrent systems or shared memory systems**
 - Composed by processes + shared objects (aka concurrent objects)
 - Processes communicate exclusively through the shared objects
 - It is a system model alternative to *message-passing*

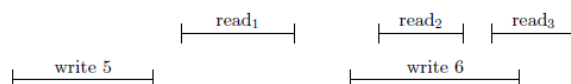


Concurrent objects

- Examples of shared objects:
 - Registers – 2 operations: read, write
 - Test&set – 1 operation: test&set
 - Compare&swap – 1 operation: compare&swap
 - FIFO queue – 2 operations: enq, deq
 - Arrays of the previous, ...
- What are they in real-life?
 - Programming language constructs for multi-threading
 - Parallel computers' shared memory and operations
 - Emulations in distributed systems

Registers

- Classification via number of readers/writers:
 - single- or multi-reader
 - single- or multi-writer
- Semantics with concurrent operations:
 - **Safe**: a read not concurrent with any write obtains the correct value (otherwise any outcome is possible)
 - **Regular**: safe + read that overlaps a write obtains either the old or new value
 - **Atomic**: safe + reads and writes behave as if they occur in some definite order



read₂ / read₃ can return 6 / 5 in a regular register (and, e.g., 42 in a safe register), but not in an atomic register

Two important questions

- How can shared memory objects be emulated on distributed systems?
 - i.e., on top of message-passing system models, e.g., the object(s) can be implemented by a (set of) server(s) accessed through RPC (e.g., using quorums for storage or consensus algorithms for implementing SMR)
- What is the “synchronization power” of each shared memory object?
 - Maurice Herlihy, “Wait-free Synchronization”, ACM TPLS, 1991
(This lecture!)

Wait-free concurrent objects

- Wait-free implementation of a concurrent data object
 - Guarantees that any process can complete any operation in a finite number of steps, regardless of the execution speeds on the other processes
- The wait-free property provides fault tolerance
 - No process can be prevented from completing an operation by undetected halting failures of other processes, or by arbitrary variations in their speed
- Fundamental problem of wait-free synchronization:
 - Given two concurrent objects X and Y, is there a wait-free implementation of X by Y?
 - E.g., Is it possible to implement a wait-free atomic register with a safe register? Is it possible to implement a FIFO queue using atomic registers?
(Yes and No, respectively)

Synchronization power of shared objects

- Given two concurrent objects X and Y , does a wait-free implementation of X by Y exist?
 - If not, **X is more “powerful” than Y**
- Herlihy defines synchronization power in terms of **consensus number**
 - Each object has an associated consensus number
 - **The consensus number for X is the largest n for which X solves consensus among n processes**
 - In an asynchronous system in which up to $n-1$ processes may crash
 - If no largest n exists, the consensus number is said to be infinite

Consensus number

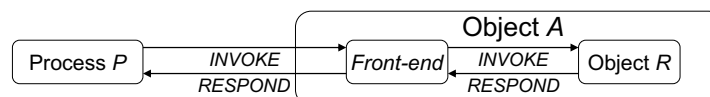
consensus number	object
1	register
2	<i>test&set, swap, fetch&add, queue, stack</i>
...	...
$2n - 2$	atomic n -register assignment
...	...
∞	<i>compare&swap, m2m move, m2m swap, augmented queue, fetch&cons, sticky byte</i>

Model

- **Concurrent system:** $[P_1, \dots, P_n; A_1, \dots, A_m]$
 - P_i : processes that invoke operations on the shared objects
 - A_i : shared objects
- **Environment:**
 - Asynchronous system
 - Up to $n-1$ processes can **crash** but objects are correct
- **Additional Assumption: Linearizability of all objects**
 - Although operations of concurrent processes may overlap, each operation appears to take effect instantaneously at some point between its invocation and response
 - In particular, operations that do not overlap take effect in their “real-time” order
 - The history “appears” sequential to each individual process
 - This apparent sequential interleaving respects the real time precedence of operations (similar to an atomic register)

Model

- An implementation of an object A is a concurrent system $\{F_1, \dots, F_n; R\}$, where the F 's are the **front-ends** and R is called the **representation object**.
 - R is the object that implements A
 - F_i is the procedure called by process P_i , to execute an operation
- We say that R implements A if there exists a wait-free implementation $\{F_1, \dots, F_n; R\}$ of A



Consensus protocol

- Consensus protocol
 - System of n processes that communicate through a set of shared objects $\{X_1, \dots, X_m\}$
 - Each process starts with an input value from some domain
 - They communicate by doing operations to the shared objects
 - They eventually agree on a common input value and halt
- A consensus protocol is required to be:
 - **Consistent**: distinct processes never decide on distinct values
 - **Wait-free**: each process decides after a finite number of steps
 - **Valid**: the common decision value is the input to some process

(these are different names for properties we know)

Consensus object

- A wait-free linearizable implementation of a consensus object is called a consensus protocol
- Consensus object (using abstract data types term.):
 - `decide(input: value) returns(value)`

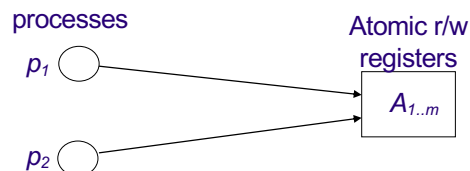
Impossibility: consensus number

- Recall the notion of **consensus number**:
 - The consensus number for X is the largest n for which X solves n -process consensus
 - If no largest n exists, the consensus number is said to be infinite.
- Theorem
 - If X has consensus number n ,
 - and Y has consensus number $m < n$,
 - then there exists no wait-free implementation of X by Y
 - in a system of more than m processes.

Impossibility: registers

- Theorem: Read/write registers have consensus number 1.
 - What does it mean to have consensus number 1?
 - Why is the theorem true?

Think about the simplest case:



Impossibility: registers

- Corollary:
 - It is impossible to construct a wait-free implementation of any object with consensus number greater than 1 using atomic read/write registers.
- Registers have no consensus power so they appear freely in the constructions that follow

Impossibility: Read-Modify-Write (RMW)

- Theorem:
 - A register with any nontrivial read-modify-write operation has a consensus number at least 2
- Examples of nontrivial rmw
 - test&set, swap, fetch&add
 - Nontrivial means the function performed is not identity

2-consensus with test&set

- test&set implementation (atomic)

```
test&set(v:value) returns(value)
  previous := v
  v := 1      //modifies local value
  return previous
end test&set
```

- 2-consensus implementation with test&set

```
shared prefer[1..2] := [1,1]; v := 0

decide(input:value) returns(value)
  prefer[p] := input
  if test&set(v) = 0
    then return prefer[p]
    else return prefer[q]
  end if
end decide
```

Theorem: test&set has consensus number 2

This is the code for process p; for process q change every p for q and vice-versa

Impossibility: test&set

- Theorem: There is no wait-free solution to three-process consensus using any combination of test&set operations
 - Generalization: No solution for any combination of read-modify-write operations that apply functions that either
 - Commute - $f_1(f_2(v)) = f_2(f_1(v))$, or
 - One function overwrites the other - $f_1(f_2(v)) = f_1(v)$
 - This includes test&set, swap and fetch&add

Compare&swap (CAS)

- CAS implementation (atomic)

```
CAS(v:value, old:value, new:value) returns(value)
  previous := v
  if previous = old then v := new endif
  return previous
end CAS
```

- Wait-free n -consensus implementation using CAS

```
shared  v :=  $\perp$ 
decide(input:value) returns(value)
  first := CAS(v,  $\perp$ , input)
  if first =  $\perp$  then return input
  else return first endif
end decide
```

- **Theorem:** A compare&swap register has infinite consensus number

Some lessons learned

- Consensus numbers:
 - Registers: 1
 - Test&set: 2
 - Compare&swap: n
- Using ... we can implement ...:
 - Compare&swap: Test&set
 - Test&set: Registers
 - Compare&swap: Registers
- Using ... we cannot implement...:
 - Test&set: compare&swap
 - Registers: Test&set
 - Registers: Compare&swap

Queues...

- **FIFO Queue**
 - enq(queue,value)
 - deq(queue):value
- **Algorithm**

```
shared prefer[1..2] := [⊥,⊥];
queue := ()

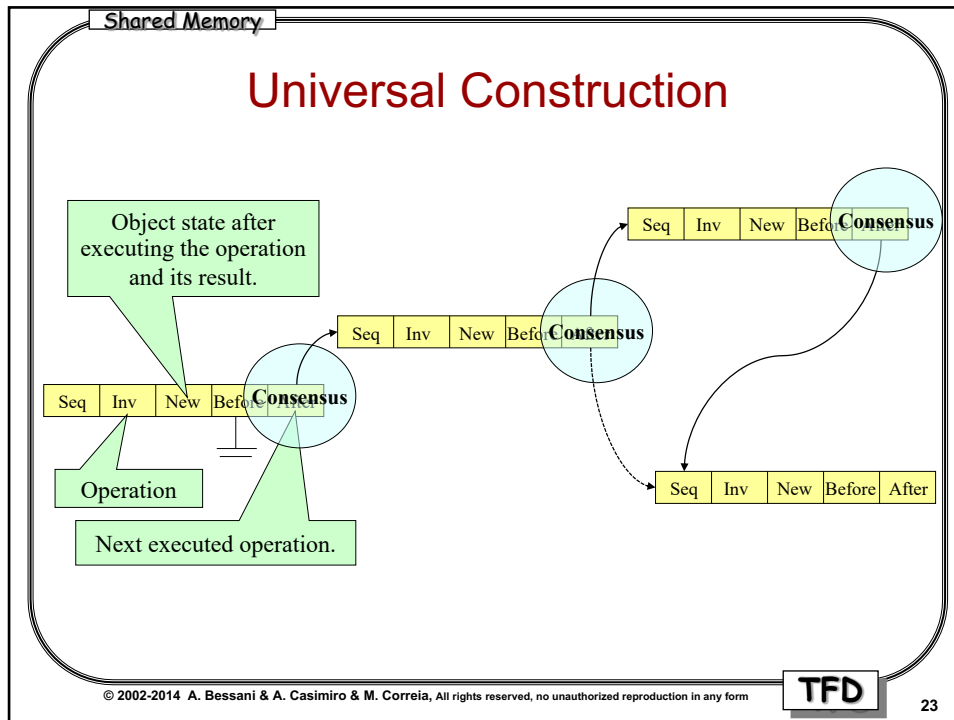
decide(input:value) returns(value)
  prefer[p] := input
  if deq(queue) = 0
    then return prefer[p]
    else return prefer[q]
  end if
end decide
```
- **Theorem: FIFO Queue**
has consensus number 2
- **Augmented Queue**
 - enq(queue,value)
 - deq(queue):value
 - peek(queue):value
- **Algorithm**

```
shared queue := ()

decide(input:value) returns(value)
  enq(queue,input)
  return peek(queue)
end decide
```
- **Theorem: Augmented Queue**
has infinite consensus number

Universality results

- An object is **universal** if it implements any other object
- Any object with consensus number **n** is universal in a system of **n** (or fewer) processes
 - E.g., compare&swap is universal
- **Basic idea for a universal construction:**
 - Represent the object as a linked list, where the sequence of cells represents the sequence of operations applied to the object
 - A process executes an operation by threading a new cell on to the end of the list
 - When the cell becomes sufficiently old, it is reclaimed and reused.
 - In other words: consensus is used to define a total order on the operations, i.e., to establish the pointer to the next element on the list



Shared Memory

Remarks

- “Wait-free Synchronization” is considered one of the most influent papers in distributed computing
 - 2003 Dijkstra Prize (distributed computing)
 - 2004 Gödel prize (theoretical computer science)
- Sort of Periodic Table (Chemistry) for shared memory distributed systems
- High impact: e.g., parallel computers now implement *compare&swap*, not *test&set*

© 2002-2014 A. Bessani & A. Casimiro & M. Correia, All rights reserved, no unauthorized reproduction in any form

TFD

24